



Mitigating Evaluative Epistemic Filtering in LLM-Driven Essay Scoring: A Culturally-Aware Multi-Perspective Scoring (CAMPS) Framework

by

Bao Nguyen

Submitted to the School of Science, Engineering and Technology
in partial fulfilment of the requirements for the degree of

Master of Artificial Intelligence

at the

RMIT UNIVERSITY

September 2026

© Bao Nguyen, 2026. All rights reserved.

Supervisor: Dr Ginel Dorleon, RMIT University

Declaration

I certify that except where due acknowledgement has been made, this thesis is my own work, that any help received in preparing this thesis has been acknowledged, and that all sources used have been indicated. I certify that the work is original and has not been submitted previously, in whole or in part, for any other academic award. Ethics procedures and guidelines have been followed (RMIT Human Research Ethics exemption, Project 31036). Any use of generative AI tools in the preparation of this thesis has been disclosed in accordance with course policy. Student name: Bao Nguyen

Student number: s4139514

Signature:

Date:

Acknowledgements

I would like to thank my supervisor, Dr Ginel Dorleon, for his guidance, direction, and foundational work on multi-perspective bias mitigation that inspired the CAMPS framework. I also thank Dr Thuy Nguyen, the course coordinator, for her support throughout the semester, and RMIT University Vietnam and the School of Science, Engineering and Technology for the resources provided during this research.

Use of generative AI. In accordance with course policy, the use of generative AI tools in this project is disclosed here. AI assistance (Claude, Anthropic) was used for software engineering support in the experiment pipeline and analysis scripts, for copy-editing and restructuring of prose drafted by the author, and for checking citations against primary sources. The research questions, framework design, experimental decisions, data collection, interpretation of results, and conclusions are the author's own; no experimental result was generated by an AI tool other than the LLM evaluators that are the object of study, and every statistic is regenerated by script from logged model outputs.

Summary

A Vietnamese student opens her essay with her grandfather's proverb, "water wears down stone", and argues as many Asian traditions teach: context first, conclusion arrived at rather than announced. A human examiner sees the skill. An AI grader trained mostly on Western text may see only a missing thesis statement and mark her down for the shape of her reasoning.

AI now grades high-stakes essays. This thesis measures that penalty and proposes CAMPS: a panel of AI judges, each reading the essay through a different cultural lens, disagreement flagging it for a teacher.

Every AI model tested scored native-English essays above upper-level Southeast Asian ones, the widest gap over a full band. Because no human had judged the two groups to be of equal quality, this is a disparity, not a proven cultural bias. The panel narrowed the gap by about a quarter on the two most disparate models, widened it on one, and more judges did not help; the disagreement signal proved too coarse, and too loosely tied to cultural disagreement, to pick out the essays needing a human. Any AI grader's gap can now be measured before it is trusted; no single cultural lens should grade alone.

Mitigating Evaluative Epistemic Filtering in LLM-Driven Essay Scoring: A Culturally-Aware Multi-Perspective Scoring (CAMPS) Framework

by

Bao Nguyen

Submitted to the School of Science, Engineering and Technology
in September 2026, in partial fulfilment of the
requirements for the degree of
Master of Artificial Intelligence

Abstract

Large language models (LLMs) increasingly score student essays, yet models trained mainly on Western text risk treating Anglo-American rhetoric as the standard and penalising other valid traditions. This thesis formalises the mechanism as *evaluative epistemic filtering* and develops CAMPS (Culturally-Aware Multi-Perspective Scoring), an inference-time framework where agents calibrated to distinct cultural readings of the rubric score each essay independently, a Bias Divergence Score (BDS) measuring disagreement. Across 200 topic-controlled ICNALE essays and four model families, every baseline model scored native-English above upper-band Southeast Asian writers (0.19–1.18 bands; all 95% intervals exclude zero). As the cohorts were not matched on judged quality, this is an unadjusted disparity, not an isolated cultural bias. A three-agent panel reduced the disparity by 23–28% on the two most disparate models, by marking both cohorts down (natives more) and by no more than its Southeast Asian lens alone; it amplified the disparity on the least disparate model. Five agents added nothing on GPT-4o and widened the gap on Claude. Reweighting showed no trade-off with band association. BDS missed its detection target (held-out sensitivity 0.06 and 0.00 at 5% false-positive rate). Culturally framed prompts change LLM scoring in model-dependent ways; cohort gaps and agent disagreement are imperfect fairness proxies.

Contents

1	Introduction	11
1.1	Problem Statement	11
1.2	Research Questions and Hypotheses	13
1.3	Theoretical Framing: Epistemic Injustice and Contrastive Rhetoric	15
1.4	Significance	16
1.5	Contributions	16
1.6	Thesis Outline	17
2	Literature Review	18
2.1	Multi-Agent and Panel-Based Approaches to Scoring	18
2.2	Inference-Time Debiasing	19
2.3	Cultural Bias in Large Language Models	20
2.4	Fairness in AES for Non-Native Writers	21
2.5	Contrastive Rhetoric and Cultural Writing Traditions	22
2.6	The Research Gap	23
3	Methodology	25
3.1	Working Theory: Evaluative Epistemic Filtering	25
3.2	The CAMPS Framework	27
3.3	Scaling the Panel: from Three to Five Agents	30
3.4	Data	32
3.5	Experimental Design	35
3.6	Alternatives Considered	36

3.7	Ethical Considerations	37
3.8	Metrics and Statistical Protocol	37
4	Results and Discussion	39
4.1	Experiment 1: The Cohort Gap at Scale (H1; H2 in part)	39
4.2	Experiment 2: Extended CAMPS Panel (H2, H4)	41
4.3	Experiment 3: BDS Validation (H3)	46
4.4	Uncertainty and Error Analysis	48
4.5	Threats to Validity	51
4.6	Discussion	53
5	Conclusion	56
5.1	Conclusions Against the Hypotheses	56
5.2	Contributions	57
5.3	Limitations	58
5.4	Future Work	59
A	CRI System Prompts	60
B	Reproducibility	63

List of Figures

1-1	Two valid shapes of argument	12
3-1	The CAMPS pipeline	27
3-2	Cohort construction and sampling flow	34
4-1	E1 cohort score gaps by model and condition	41
4-2	Model-dependent panel effect	42
4-3	Lens-level cohort profiles (radar)	44
4-4	E2 ablation: gap by panel configuration	45
4-5	E2 gap versus band association under reweighting (GPT-4o)	46
4-6	E3: BDS against human-rater disagreement	47
4-7	Baseline gap by SEA proficiency band	49
4-8	One essay, every lens	51

List of Tables

1.1	Mapping of research questions to hypotheses	14
1.2	Key terms used in this thesis	14
3.1	Capability comparison of related frameworks	30
3.2	The five-lens CAMPS panel	31
3.3	ICNALE modules used in this thesis	33
3.4	Overview of Experiments E1–E3	35
4.1	E1: cohort score gap by model	42
4.2	E2: ablation over panel size and adjudicator	45
4.3	E3: BDS detection performance	47
5.1	Verdicts on hypotheses H1–H4	57

Abbreviations

ASAP	Automated Student Assessment Prize (essay dataset)
AUC	Area Under the ROC Curve
AES	Automated Essay Scoring
AI	Artificial Intelligence
API	Application Programming Interface
BDS	Bias Divergence Score
CAMPS	Culturally-Aware Multi-Perspective Scoring
CEFR	Common European Framework of Reference for Languages
CI	Confidence Interval
CRI	Cultural Rubric Interpretation
ENS	English Native Speaker (ICNALE cohort)
FPR	False Positive Rate
GRA	Global Rating Archives (ICNALE module)
ICNALE	International Corpus Network of Asian Learners of English
IELTS	International English Language Testing System
L1	First (native) language
LLM	Large Language Model
QWK	Quadratic Weighted Kappa
ROC	Receiver Operating Characteristic
SD	Standard Deviation
SEA	Southeast Asia(n)
TOEFL	Test of English as a Foreign Language
WE	Written Essays (ICNALE module)
WEIRD	Western, Educated, Industrialised, Rich, Democratic
WEP	Written Essays Plus (ICNALE module)

Chapter 1

Introduction

1.1 Problem Statement

Automated Essay Scoring (AES) has entered a new era. Where earlier systems counted surface features such as word length, sentence complexity, and spelling, large language models (LLMs) now evaluate what previous systems could not: the quality of a student’s argument. Frontier models are increasingly deployed to score high-stakes writing at scale, promising consistent, rapid, and inexpensive assessment. Yet a model that judges *how well a student argues* must hold some standard of what good argumentation looks like, and that standard is not neutral. LLMs derive their evaluative expectations from training corpora that overwhelmingly represent Western, Educated, Industrialised, Rich, and Democratic (WEIRD) cultural contexts; empirical work has shown that widely used models express cultural values closely aligned with English-speaking and Protestant-European countries [Tao et al., 2024].

Consider an illustration. A Vietnamese student opens an argumentative essay with her grandfather’s proverb, *nuoc chay da mon* (water wears down stone). She then develops her position the way many Asian academic writing traditions teach: context first, evidence accumulated gradually, the conclusion arrived at rather than announced. A trained human examiner can recognise the coherence of this inductive structure. An LLM evaluator calibrated by its training distribution to expect a thesis statement in the opening paragraph may instead record “lacks a clear thesis” and “argument feels indirect”. It may also penalise the

hedging forms that signal politeness in her tradition as weak reasoning. If so, she loses marks not because her English is poor, but because the shape of her reasoning is culturally unfamiliar to the model. Figure 1-1 sketches the two shapes.




Figure 1-1: Two valid shapes of argument. The writer-responsible, thesis-first pattern is plausibly dominant in the text LLM evaluators learn from, given their measured alignment with Western values [Tao et al., 2024]; the reader-responsible, context-first pattern common in many Asian academic traditions [Hinds, 1987, Connor, 2011] is equally valid but under-represented.

This thesis names this mechanism *evaluative epistemic filtering*: when an LLM evaluates writing, it applies an implicit cultural filter derived from its training distribution, systematically discounting rhetorical strategies that deviate from Anglo-American conventions regardless of their argumentative merit. A pilot study in the preceding phase of this research (Part A) reported scoring disparities between Southeast Asian (SEA) and native-English cohorts on a 30-essay sample across two frontier models, and that a culturally-diversified scoring panel reduced them. The pilot's corpus did not support a controlled comparison, so Part B rebuilds the data (Section 3.4); its findings motivated, but do not establish, the claims this thesis investigates at scale.

The problem this thesis addresses can therefore be stated plainly. If LLM-based essay scoring applies culturally partial standards, it does so invisibly and at scale, to students whose academic futures depend on the result. Scoring disparities against non-native writers are documented, but no validated framework exists to mitigate them in evaluative settings, and no calibrated mechanism exists to flag the essays affected. This thesis measures the native–SEA scoring disparity across models, builds and tests such a framework, and tests its flagging mechanism against human ground truth. Throughout, *disparity* or *gap* names what is measured (a difference in mean score between cohorts) and *filtering* names the mechanism hypothesised to produce part of it (Table 1.2).

1.2 Research Questions and Hypotheses

Two research questions drive this thesis:

- RQ1.** To what extent does evaluative epistemic filtering manifest as measurable scoring disparities against Southeast Asian rhetorical patterns in LLM-driven essay scoring, and does it generalise across model families?
- RQ2.** Can a culturally-aware multi-agent scoring architecture mitigate these disparities at inference time, without retraining or access to model weights, while reliably flagging for human review the essays it cannot score with confidence?

These questions are operationalised as four testable hypotheses:

- **H1 (Generalisation).** A native—SEA cohort score gap, consistent with evaluative epistemic filtering, has a 95% confidence interval excluding zero on at least four LLMs, including an open-weight model.
- **H2 (Mitigation).** A five-agent culturally-calibrated panel with a meta-adjudicator reduces the cohort score gap more than the three-agent pilot configuration, without degrading agreement with the corpus’s proficiency bands.
- **H3 (Detection).** A Bias Divergence Score threshold calibrated by receiver operating characteristic (ROC) analysis attains at least 90% sensitivity at no more than a 5% false-positive rate against multi-rater human disagreement on the 140-essay Global Rating Archives (GRA) validation set.
- **H4 (Trade-off boundary).** Consensus-weight allocation exhibits an identifiable boundary beyond which cohort-fairness gains degrade scoring accuracy against graded ground truth.

Graded ground truth does not exist for the cohort essays, so H4 is examined on a proxy, association with the corpus’s proficiency bands (Section 3.8), and is reported as exploratory.

The research questions and hypotheses relate as follows. RQ1 is a question about the *existence and generality* of the problem; it is operationalised by H1, which makes it falsifiable

by fixing the scope (four model families) and the test (a confidence interval on the cohort scoring gap that excludes zero). RQ2 is a question about the *solution*, and it decomposes into three parts: whether the mitigation works (H2), whether its failures can be detected (H3), and what the mitigation costs (H4). Table 1.1 summarises the mapping and where each hypothesis is tested.

Table 1.1: How the research questions decompose into hypotheses, and where each is tested.

RQ	Hypothesis	Establishes	Tested in
RQ1	H1	A cohort gap exists and generalises across models	E1
RQ2	H2	The panel narrows the gap	E1–E2
RQ2	H3	BDS flags contested essays	E3
RQ2	H4	What narrowing the gap costs (proxy: band association)	E2

Chapter 5 returns to each hypothesis explicitly. Because several terms recur throughout, Table 1.2 defines them once in plain language.

Table 1.2: Key terms used in this thesis, in plain language.

Term	Meaning in this thesis
Evaluative epistemic filtering	The tendency of an AI grader to mark down writing whose rhetorical form is unfamiliar from its training data, regardless of the argument’s quality.
Lens	One AI evaluator given a documented cultural interpretation of the same scoring rubric (Western, Southeast Asian, East Asian, Mekong, or Neutral).
Panel, consensus	Several lenses scoring the same essay independently; their weighted average is the consensus score.
Bias Divergence Score (BDS)	How much the lenses disagree on one essay; high values are meant to flag the essay for a human.
Cohort, cohort gap	A group of essays by writer background (Southeast Asian learners or native-English writers); the gap is the difference in mean score between the two cohorts. It is an <i>unadjusted disparity</i> : the cohorts were not matched on human-judged quality, so the gap may contain both filtering and real quality differences.
Contested essay	In E3, an essay in the top quartile of disagreement among 80 human raters. This is a proxy for the cultural contestation BDS is meant to detect, not the same thing.
Band	The 1–9 score an evaluator assigns (integer, with occasional half-bands); the corpus’s own proficiency labels are CEFR bands (B1_2, B2+).

1.3 Theoretical Framing: Epistemic Injustice and Contrastive Rhetoric

Two established bodies of theory ground this thesis. The first is Fricker [2007]’s account of *testimonial injustice*: a speaker is wronged in their capacity as a knower when prejudice in the hearer deflates the credibility of their testimony. Evaluative epistemic filtering can be understood as testimonial injustice automated and scaled: the student’s argument is discounted not on its merits but because its form does not match the evaluator’s culturally conditioned expectations. Fricker’s framework also cautions against a naive remedy: no evaluator, human or artificial, judges from nowhere. This thesis therefore treats a “culturally neutral” evaluator as a regulatory ideal rather than an achievable design goal, and pursues *plurality* of perspectives rather than pretended neutrality.

The second grounding is contrastive rhetoric. Since Kaplan [1966] first observed that rhetorical organisation is culturally learned, research summarised by Connor [2011] has documented systematic differences between writing traditions. These include the distinction between writer-responsible and reader-responsible coherence, and the validity of inductive, context-first argument structures common in Asian academic writing. This literature supplies the scholarly basis for the cultural rubric interpretations at the centre of the proposed framework, distinguishing them from ad-hoc cultural personas.

The mitigation strategy itself extends the formal multi-perspective bias-reduction framework of Dorleon and Shujaa [2026], who proved that, under ideal conditions, generating text from multiple cultural perspectives with bias filtering can eliminate stereotypical content and perspective omission in *generative* tasks. Companion empirical work documented pervasive cultural bias in LLM treatments of Southeast Asian content [Shujaa et al., 2026]. This thesis asks whether those principles transfer from generation to *evaluation*. Independent support for the core mechanism comes from the evaluation literature: panels of diverse LLM judges have been shown to outperform and de-bias single large judges [Verga et al., 2024], though with diversity defined by model provider rather than by culture. The thesis, in short, tests whether the jury effect holds when the axis of diversity is cultural.

Finally, the fairness–accuracy trade-off at the centre of H4 has a direct lineage in fairness-

aware machine learning. Dorleon et al. [2022] formulated feature selection as an explicit trade-off between fairness and prediction performance. FAPFID [Dorleon et al., 2023] then showed that fairness for unprivileged groups can be improved without sacrificing overall performance, through balanced representation. H4 asks whether an analogous frontier exists in evaluative scoring, and where it lies.

1.4 Significance

The significance of this work is threefold. First, equity: as LLM-based assessment becomes a gatekeeper to scholarships, university admission, and migration pathways, a systematic penalty against non-Western rhetorical traditions would constitute discrimination embedded in infrastructure, invisible to its subjects and deniable by its operators. Quantifying that penalty rigorously is a precondition for governing it. Second, method: existing multi-agent AES research improves accuracy without cultural awareness [e.g., Su et al., 2025], while existing debiasing research targets generative outputs rather than evaluative judgments; a validated framework for *evaluative* fairness would fill a documented gap at their intersection. Third, deployability: because the proposed framework operates entirely at inference time through commercial APIs, any institution can apply it without retraining models. This practical property distinguishes it from training-time approaches and matches the constraints of real educational deployments, particularly in the Southeast Asian context where this research is situated.

1.5 Contributions

This thesis makes five contributions:

1. **Conceptual:** it formalises evaluative epistemic filtering as a distinct bias mechanism in LLM-driven scoring, theoretically grounded in epistemic injustice and contrastive rhetoric.
2. **Architectural:** it develops CAMPS (Culturally-Aware Multi-Perspective Scoring),

a culturally-calibrated scoring panel with a divergence-based review trigger, extending multi-perspective bias reduction from generation to evaluation, and tests whether scaling it to five lenses and an adjudicator helps.

3. **Empirical:** it evaluates the framework on 200 topic-controlled essays from the ICNALE corpus [Ishikawa, 2023] across four frontier and open-weight models, establishing the disparity at scale on a corpus suited to the comparison and characterising where the framework changes it.
4. **Methodological:** it introduces ROC calibration of the Bias Divergence Score against multi-rater human disagreement, replacing the pilot’s hand-set threshold.
5. **Practical:** it maps the gap against a proficiency-band association proxy under every panel weighting, giving deploying institutions a re-runnable audit and an explicit account of what narrowing the gap did and did not cost on this evidence.

1.6 Thesis Outline

The thesis follows the research questions in order. Chapter 2 reviews the literature that frames RQ1 and RQ2 and locates the gap this thesis occupies. Chapter 3 develops the CAMPS framework and the method: the data, the three experiments (E1 for H1 and part of H2, E2 for H2 and H4, E3 for H3), the statistical protocol, and the alternatives considered. Chapter 4 reports the experiments in that order and discusses what the results mean, why they arose, how they compare with prior work, and where they are uncertain. Chapter 5 answers each hypothesis, states the contributions and limitations, and outlines future work.

Chapter 2

Literature Review

The literature relevant to this thesis was searched between July and September 2026 using Google Scholar, arXiv, the ACL Anthology, and the Springer and IEEE digital libraries. The search combined the terms “automated essay scoring”, “multi-agent”, “LLM-as-a-judge”, “cultural bias”, “debiasing”, “fairness”, “non-native writers”, and “contrastive rhetoric”. It was supplemented by backward and forward citation chasing from key papers. Every reference cited in this thesis was verified against its primary source. The review is organised thematically: multi-agent and panel-based scoring (Section 2.1), inference-time debiasing (Section 2.2), cultural bias in LLMs (Section 2.3), fairness in AES for non-native writers (Section 2.4), and the contrastive-rhetoric foundations (Section 2.5), before Section 2.6 locates the gap this thesis addresses.

2.1 Multi-Agent and Panel-Based Approaches to Scoring

Two research communities have converged on the same architectural intuition, that several evaluators are better than one, for different reasons.

In automated essay scoring, multi-agent designs pursue *accuracy*. CAFES [Su et al., 2025] orchestrates scorer, feedback-manager, and reflective-reviser agents (average 21% improvement in Quadratic Weighted Kappa, QWK, over direct prompting). Roundtable Essay Scoring [RES; Jang et al., 2025] consolidates trait-based evaluator agents through simulated discussion (up to 34.86% QWK improvement on the ASAP benchmark, zero-shot). Related

designs add debate with retrieval [Keramati et al., 2026] and component recognition [Wang et al., 2025]. What unites this family is the axis along which agents differ: *function* or *trait*. None varies the cultural interpretation of the rubric itself, and none evaluates fairness across writer cohorts. Improvements are reported on aggregate agreement. This leaves open the possibility that a more accurate panel is uniformly accurate for some writers and uniformly harsh for others.

In the LLM-evaluation community, panel designs pursue *bias reduction*. Verga et al. [2024] replaced a single large judge model with a Panel of LLM evaluators (PoLL) drawn from three different model families. They found that the panel not only correlated better with human judgement across six datasets but also reduced intra-model self-preference bias, at roughly a seventh of the cost. This is, to date, the clearest empirical demonstration that judge *diversity* counteracts judge *bias*. The diversity axis, however, is model provider, chosen to decorrelate model-specific quirks rather than to represent any substantive difference in evaluative standards. The present thesis can be read as a targeted transplant. It takes the panel-of-judges mechanism whose bias-reducing effect Verga et al. [2024] established, and re-axes the diversity from provider to *cultural rubric interpretation*. Here the bias to be countered is cultural rather than idiosyncratic.

2.2 Inference-Time Debiasing

A second literature attacks bias at inference time, without retraining. BiasFilter [Cheng et al., 2025] periodically evaluates candidate continuations during generation against a fairness reward model and discards low-reward segments, enforcing fairness on *generated text* for both open-source and API-based models. The framework demonstrates that meaningful debiasing is achievable purely at inference time, a design constraint this thesis shares. But its unit of intervention is the output token stream, which has no analogue when the model’s output is a single scalar score.

Closest to the present work is Multi-Agent Cultural Debate (MACD) [Tan et al., 2026], which assigns LLM agents explicit cultural personas and orchestrates deliberation. MACD substantially raised the rate of unbiased responses over single-model baselines, and its cen-

tral finding, that functionally-rolled agent frameworks lack the explicit cultural representation cross-cultural fairness needs, independently corroborates the design principle behind CAMPS. The two frameworks nonetheless answer different questions. MACD debiases *what a model says* in open-ended generation. CAMPS debiases *how a model judges*: the deliverable is a score whose fairness can be measured cohort by cohort, and disagreement between cultural agents is exploited as a detection signal rather than dissolved through debate.

A caution for all multi-agent designs comes from the MALIBU benchmark [Mirza et al., 2025], which shows that LLM-based multi-agent systems can implicitly *reinforce* social biases, and that persona-conditioned judging can over-correct. For this thesis, MALIBU motivates two design decisions: consensus weights are calibrated empirically rather than assumed benign (Hypothesis H4 measures the over-correction frontier directly), and panel behaviour is validated against multi-rater human ground truth rather than taken on faith.

A separate line of work mitigates bias inside the model rather than around it. Self-debiasing exploits a model’s own ability to recognise undesirable properties of its output, suppressing them at decoding time [Schick et al., 2021], and FairSteer intervenes on hidden-layer activations with dynamically applied steering vectors at inference time [Li et al., 2025]. Both are effective in their settings, but both are white-box methods: they require access to output distributions or internal activations that commercial scoring deployments accessed through an API do not expose. CAMPS was deliberately designed to operate at the level of prompts and outputs alone, trading the precision of internal intervention for applicability to any evaluator, including closed models.

2.3 Cultural Bias in Large Language Models

That LLMs carry a Western-centric value orientation is well documented, although whether it extends to *evaluative* judgement is exactly what this thesis tests. Tao et al. [2024] compared five OpenAI model generations against nationally representative values-survey data for 107 countries. They found that model responses consistently resemble those of English-speaking and Protestant-European populations. This is the WEIRD alignment that this the-

sis takes as the distal cause of evaluative epistemic filtering.

The proximal framework for this thesis comes from the supervisor’s research programme. Dorleon and Shujaa [2026] formally defined bias categories for generative outputs (stereotyping, omission, ethnocentrism, simplification) and proposed multi-perspective generation with bias detection, proving that under ideal conditions the method eliminates stereotypical content and perspective omission. The companion study [Shujaa et al., 2026] evaluated seven state-of-the-art LLMs on prompts embedding culturally sensitive assumptions about Vietnam, Myanmar, and Nepal, documenting pervasive stereotyping, linguistic ethnocentrism, and epistemic bias, with substantial variation across models. The same group’s work on framing-bias detection [Shujaa and Dorleon, 2025] reflects the detection-oriented strand of this programme that the Bias Divergence Score continues: bias treated not only as something to remove but as something to *measure and flag*.

Two further studies shaped how cultural calibration is implemented in this thesis. SEA-HELM, a community-built evaluation suite for Southeast Asian languages, argues that benchmarks produced by machine translation or synthetic generation without community input lose cultural authenticity [Susanto et al., 2025]. The same reasoning led this thesis to ground its cultural lenses in the contrastive rhetoric literature and in a corpus built in the region, rather than in ad hoc persona descriptions. More cautionary still, Khan et al. [2025] showed that survey-based measures of LLM cultural alignment are unstable across presentation formats and that prompting a model to represent a specific culture is not reliably steerable. For CAMPS this is a design constraint, not a refutation: lens prompts are fixed verbatim (Appendix A), decoding is deterministic, and the panel’s behaviour is validated empirically against human raters (E3) rather than assumed correct because the prompt requested it.

2.4 Fairness in AES for Non-Native Writers

Concerns about equitable automated scoring of non-native writing predate the LLM era. Vajjala [2018] showed on the TOEFL11 and FCE corpora that the linguistic features most predictive of essay quality differ across L1 populations, evidence that a single scoring model implicitly encodes population-specific assumptions about what good writing looks like.

What has changed with LLM-based scoring is the depth at which such assumptions operate: where feature-based systems encoded them in engineered predictors that could be inspected, LLM evaluators encode them in an opaque prior over rhetorical form [cf. Gallegos et al., 2024].

The LLM-era evidence is accumulating quickly. Pack et al. [2024] evaluated four prominent LLMs (PaLM 2, Claude 2, GPT-3.5, GPT-4) as scorers of English-learner placement essays, finding usable reliability for the strongest model but meaningful variation across models and administrations. Evaluative behaviour is therefore model-specific, which motivates testing H1 across at least four model families rather than generalising from one. Most directly relevant, Gayed [2026] fine-tuned an open-weight model (Gemma-3-27B) for essay scoring and evaluated it on the full TOEFL11 corpus of 12,100 essays across eleven L1 backgrounds, explicitly analysing scoring effects by first language. The study documents L1-linked scoring variation on exactly the corpus family and task that concern this thesis. Like the broader AES fairness literature it extends, however, it stops at *measurement*: no mitigation framework is proposed, and the native/non-native design this thesis requires is out of scope because TOEFL11 contains no native-English cohort. Related recent work has begun to probe mechanism as well as outcome. Yang et al. [2025] showed that prompt-based LLM scorers can infer writer demographics, including English-learner status, from essay text alone. They also showed that scoring behaviour shifts once such a status is recognised.

Taken together, this literature establishes the disparity, its model-dependence, and hints at its mechanism, while leaving the mitigation side of the problem essentially untouched. That asymmetry between diagnosis and treatment is the practical motivation for the framework developed in Chapter 3.

2.5 Contrastive Rhetoric and Cultural Writing Traditions

The claim that rhetorical organisation is culturally learned, and that an evaluator calibrated to one tradition will therefore systematically misread another, has a research lineage of its own. Kaplan [1966] first argued that paragraph organisation reflects culturally specific “thought patterns”, contrasting the linear development expected in English academic prose

with the patterns he observed in other traditions. The framework was refined substantially in later work: Hinds [1987] distinguished *writer-responsible* languages, in which the burden of clarity falls on the writer through explicit signposting, from *reader-responsible* languages, in which coherence is expected to be inferred by the reader. Under this typology, the “missing thesis statement” an LLM penalises may simply be reader-responsible coherence operating as designed. Connor [2011] consolidated this field as intercultural rhetoric, correcting early essentialist readings: rhetorical traditions are conventions distributed across writing communities, not fixed properties of individuals. This thesis adopts that distinction explicitly, treating L1 cohorts as proxies for rhetorical tradition rather than for identity.

For Southeast Asian learner writing specifically, the empirical base is anchored by the ICNALE project [Ishikawa, 2023], whose topic-controlled design allows rhetorical variation to be observed under matched task conditions across ten Asian countries and a native-English reference cohort. This literature performs two duties in the present thesis. First, it supplies the content of the cultural rubric interpretations in Chapter 3, where each Cultural Rubric Interpretation (CRI) prompt is grounded in documented rhetorical conventions rather than improvised cultural personas. Second, it supplies the defence against the objection that culturally-calibrated agents merely role-play stereotypes. The calibrations are drawn from a half-century of contrastive-rhetoric scholarship, and their behaviour is tested empirically in Chapter 4.

2.6 The Research Gap

Four literatures converge on the problem this thesis addresses, and each stops short of it in a characteristic way. The multi-agent scoring literature (Section 2.1) established that panels of evaluators outperform single judges (CAFES and RES on accuracy, PoLL on bias). But it defines panel diversity by function, trait, or model provider, never by cultural interpretation of the rubric, and it does not measure fairness across writer cohorts. The inference-time debiasing literature (Section 2.2) established that bias can be mitigated without retraining, and MACD specifically established that *explicit cultural representation* in agent frameworks matters. But its object is generated text, not evaluative judgement,

and it dissolves inter-agent disagreement rather than exploiting it. The cultural-bias literature (Section 2.3) established the WEIRD disposition of current models and, in Dorleon and Shujaa [2026], provided a formal multi-perspective mitigation framework with provable guarantees, for generation only. Meanwhile, the AES fairness literature (Section 2.4) continues to document L1-linked scoring disparities without offering a validated mitigation framework.

Within the scope of the searches described above, no published work was found at the intersection: culturally-diverse rubric interpretation, applied to evaluative scoring, with inter-agent divergence used as a calibrated trigger for human review and tested against multi-rater human ground truth. That intersection is where CAMPS operates.

Chapter summary

Two literatures converge on the idea that several evaluators are better than one, but neither varies the *cultural* reading of the rubric, and the fairness literature documents L1-linked scoring disparities without a validated remedy. Chapter 3 builds the framework that fills this gap and the experiments that test it.

Chapter 3

Methodology

This chapter sets out what was done and why. It states the working theory (Section 3.1) and specifies the CAMPS framework (Sections 3.2 and 3.3). It then describes the data (Section 3.4) and the three experiments (Section 3.5), records the alternatives considered and the ethical constraints (Sections 3.6 and 3.7), and fixes the metrics and statistical rules before any result is examined (Section 3.8).

3.1 Working Theory: Evaluative Epistemic Filtering

The framework developed in this chapter rests on a working theory of why LLM evaluators systematically under-score culturally unfamiliar writing. This section states that theory precisely, identifies the conditions that sustain the phenomenon, and derives the observable signatures that the experiments in Chapter 4 test.

Definition 3.1 (Evaluative epistemic filtering). *Let E be an essay, $f_{\text{human}}(E)$ the score assigned by a trained human examiner under a rubric R , and $f_{\text{LLM}}(E \mid \theta)$ the score assigned by an LLM evaluator whose parameters θ were estimated from a training distribution $\mathcal{D}$, and let $c(E)$ denote the rhetorical tradition of the essay. Evaluative epistemic filtering occurs when the evaluator’s deviation from human judgement differs between traditions,*

$$\mathbb{E}[f_{\text{LLM}} - f_{\text{human}} \mid c(E) = c_1] \neq \mathbb{E}[f_{\text{LLM}} - f_{\text{human}} \mid c(E) = c_2], \quad (3.1)$$

with essays whose rhetorical form is under-represented in $\mathcal{D}$ receiving the more deflated scores relative to human judgement, independent of argumentative merit. The between-tradition contrast is essential: an evaluator that scores every tradition one band too low is miscalibrated but does not filter.

The definition makes two commitments. First, the filter is a property of the *evaluator*, not of the essay: the same essay receives different treatment under evaluators with different training distributions. Second, the deviation is conditional on rhetorical tradition rather than on language proficiency. This is what distinguishes evaluative epistemic filtering from the legitimate scoring of weaker English. It is what the band-restricted cohort design of Section 3.4 is built to approach, and what a supplementary human-anchored check (Section 4.4) tests where human ratings exist.

Four conditions jointly sustain the filter. First, training corpora over-represent WEIRD cultural contexts, whose value orientation models demonstrably absorb [Tao et al., 2024]. Second, the standards LLM evaluators apply are opaque, distributed across parameters rather than inspectable as engineered features [Vajjala, 2018, Gallegos et al., 2024]. Third, diagnostic effort has concentrated on generative outputs, leaving the bias categories of Dorleón and Shujaa [2026] without an evaluative counterpart. Fourth, the observable symptom, L1-linked scoring disparity, is documented [Gayed, 2026] but unconnected to any mitigation mechanism.

The theory yields three testable signatures, mapping onto the hypotheses of Chapter 1. The first is a systematic scoring gap between an upper-band learner cohort and the native reference, consistent with filtering though not separable by this design from residual quality difference (H1). The second is shrinkage of that gap under a panel whose members embody different rubric interpretations, on the premise that their individual filters differ (H2); independent calls do not by themselves make the lenses' errors statistically independent, since they share a base model. The third is measurable panel disagreement on precisely the essays where filtering operates, making divergence itself a detection signal (H3). The fairness–accuracy boundary (H4) follows from the consensus mechanism below.

3.2 The CAMPS Framework

CAMPS (Culturally-Aware Multi-Perspective Scoring) operationalises a single design principle: if no evaluator judges from nowhere (Section 1.3), fairness is pursued through a *plurality* of explicitly declared perspectives rather than a pretended neutrality. In practice this means five cultural rubric interpretations (Table 3.2), each declared in an inspectable system prompt rather than hidden in model weights. The framework operates entirely at inference time, requires no retraining or weight access, and treats every underlying LLM as a black box reachable through an API. Figure 3-1 shows the pipeline; the five components are specified below.




Figure 3-1: The CAMPS pipeline. Each essay is scored in parallel by k *culturally-calibrated* agents: LLM evaluators whose system prompts embed a documented cultural interpretation of the scoring rubric (Component 1). The consensus engine aggregates the scores. The Bias Divergence Score routes contested essays to human review, whose outcomes feed the calibration loop. CAMPS components are shaded blue; inputs and outcomes are unshaded.

Component 1: Cultural Rubric Interpreter (CRI)

Given the scoring rubric R , the CRI produces k culturally-adapted interpretations $\{r_1, \dots, r_k\}$ through system-level prompting. Each interpretation is grounded in the contrastive-rhetoric literature reviewed in Section 2.5 rather than in improvised personas. The Southeast Asian interpretation, for example, instructs the agent to recognise inductive thesis development and reader-responsible coherence [Hinds, 1987] as valid realisations of the rubric’s coherence criterion. The neutral interpretation is retained as a regulatory ideal in Fricker’s sense [Fricker, 2007], not as an achievable standpoint. The full prompt texts are reproduced verbatim in Appendix A.

Component 2: Evaluator agent pool

Each interpretation is bound to an independent evaluator agent, and each agent scores the essay without access to the other agents’ outputs:

$$s_i = A_i(E, r_i), \quad i = 1, \dots, k, \quad (3.2)$$

with temperature fixed at $t = 0.0$ to minimise generation variability and structured (JSON) output; every response is archived, so the analysis is exactly reproducible from the archived outputs even though provider APIs do not guarantee bit-identical regeneration. Independence at this stage is deliberate: the MALIBU results [Mirza et al., 2025] caution that interacting agents can amplify rather than cancel bias, so CAMPS confines interaction to the aggregation stage, where it is mathematically controlled.

Component 3: Consensus aggregation

The panel’s scores are combined as a weighted mean,

$$S^*(E) = \frac{\sum_{i=1}^k w_i s_i(E)}{\sum_{i=1}^k w_i}, \quad w_i \geq 0, \quad (3.3)$$

with weights initialised equal. Because S^* is linear in the agents' scores, the consensus cohort gap is the same weighted combination of the individual lenses' cohort gaps. Two consequences follow. Reweighting can trade the gap against agreement with graded ground truth, which is the boundary hypothesised in H4. But when every lens's gap has the same sign, the gap-minimising weighting is always the single lens with the smallest gap; a minimum-gap result therefore reflects the aggregation rule, not evidence that combining perspectives improves fairness. Section 3.5 describes how the weight space is swept, and Section 4.2 reports the individual lenses alongside the panel for exactly this reason.

Component 4: Bias Divergence Score (BDS)

The framework's detection signal is the dispersion of the panel:

$$\text{BDS}(E) = \sqrt{\frac{1}{k} \sum_{i=1}^k (s_i(E) - S^*(E))^2}. \quad (3.4)$$

BDS is the dispersion of the prompt-conditioned scores around the consensus (computed with equal weights in all experiments). A low BDS indicates agreement among the lenses; agreement alone does not establish that the score is accurate or fair, since all lenses may share a bias. A high BDS means the lenses disagree, which the working theory proposes as the signature of an essay whose score depends on culturally contingent expectations; whether that interpretation holds is what E3 tests. Under the framework, such an essay is flagged for human review. Part B replaces the pilot's hand-set threshold ($\tau = 0.75$) with one calibrated by ROC analysis against multi-rater human ground truth (H3; Section 3.5).

Component 5: Calibration loop

Outcomes of human review feed back into the framework in two ways: the consensus weights w_i are re-optimised to maximise agreement with human scores subject to a cohort-fairness constraint, and the threshold τ is re-estimated as review outcomes accumulate. The loop closes the system without touching model weights; everything learned lives in w and τ . The loop is specified as part of the framework but not exercised in this thesis, whose

experiments evaluate a single calibration pass (E2’s weight sweep and E3’s threshold calibration).

The five components divide the framework’s labour: the first two create and apply the plurality of perspectives, the third converts plurality into a single defensible score, and the last two turn residual disagreement into detection and adaptation. The design intends the panel’s fairness properties to be a property of the assembly rather than of any one lens; whether they are is tested directly in Section 4.2. Table 3.1 positions CAMPS against the frameworks reviewed in Chapter 2 along the five capabilities the design combines.

Table 3.1: Architectural features compared (what each framework is designed to do, not what has been validated). *Inf.* = inference-time only; *Cult.* = explicit cultural perspectives; *Eval.* = evaluative (scoring) task; *Div.* = divergence-based detection; *HITL* = a threshold-calibrated human-in-the-loop trigger. For CAMPS the detection and trigger features are implemented and evaluated in E3; their operational validity is a finding of Chapter 4, not a premise.

Framework	Inf.	Cult.	Eval.	Div.	HITL
CultureLLM [Li et al., 2024]	×	✓	×	×	×
BiasFilter [Cheng et al., 2025]	✓	×	×	×	×
CAFES [Su et al., 2025]	✓	×	✓	×	×
RES [Jang et al., 2025]	✓	×	✓	×	×
PoLL [Verga et al., 2024]	✓	×	✓	×	×
MACD [Tan et al., 2026]	✓	✓	×	×	×
Dorleon & Shujaa [2026]	✓	✓	×	×	×
CAMPS (this thesis)	✓	✓	✓	✓	✓

Relative to the Part A pilot, the framework specified here differs in three respects, each developed in the following sections. The panel is scaled from three to five cultural interpretations with a meta-adjudicator (Section 3.3). The corpus and cohort design are rebuilt on data that supports the comparison (Section 3.4). Both the consensus weights and the BDS threshold are calibrated empirically rather than hand-set (Section 3.5).

3.3 Scaling the Panel: from Three to Five Agents

The pilot panel comprised three rubric interpretations: Western, Southeast Asian, and a neutral regulatory ideal. Scaling to five requires a principled answer to *which* lenses to add.

A candidate lens must satisfy three criteria. It must be (i) *documented* in the contrastive-rhetoric literature, so the CRI prompt encodes scholarship rather than stereotype; (ii) *validatable* against a corpus cohort from the corresponding tradition; and (iii) *distinct*, expected to diverge from the existing panel on identifiable features. These criteria select an *East Asian* lens, grounded in reader-responsible coherence [Hinds, 1987] and validatable against ICNALE’s East Asian cohorts, and a *Mekong/Vietnamese* lens, motivated by the region this research serves and validatable against the WEP cohorts. A Middle Eastern lens fails the validatable criterion for this corpus and is deferred. The Mekong lens is a partial exception to criterion (i): the contrastive-rhetoric literature on Mekong-region academic writing is thin, and the lens is admitted on regional relevance (Table 3.2). Separate validation of the two added lenses was outside the scope of E1–E3, which test the panel as a whole. Table 3.2 summarises the panel.

Table 3.2: The five-lens panel: rhetorical grounding and the corpus cohort against which each lens’s behaviour can be validated.

Lens	Rhetorical grounding	Validating cohort (available)
CRI _{Western}	Writer-responsible coherence; explicit thesis-first linear development [Kaplan, 1966, Connor, 2011]	ICNALE ENS
CRI _{SEA}	Inductive, context-first development; collectivist framing [Connor, 2011, Ishikawa, 2023]	ICNALE THA, IDN, PHL
CRI _{EastAsian}	Reader-responsible coherence; delayed thesis; <i>ki-sho-ten-ketsu</i> -style organisation [Hinds, 1987]	ICNALE CHN, JPN, KOR, TWN
CRI _{Mekong}	Proverb- and narrative-anchored argumentation; scholarly motivation from the ICNALE Guide and the supervisor’s cultural-bias work [Ishikawa, 2023, Shujaa et al., 2026], not an empirical study of Mekong learner rhetoric	ICNALE WEP VNM, KHM, LAO, MMR
CRI _{Neutral}	Regulatory ideal: argumentative merit irrespective of structural ordering [Fricker, 2007]	(all cohorts)

The second architectural addition is a *meta-adjudicator*: a further LLM instance that receives the five agents’ scores and their short written rationales, and produces a reasoned

consensus score,

$$S^{\text{adj}}(E) = M(s_1, \rho_1, \dots, s_k, \rho_k), \quad (3.5)$$

where ρ_i is agent i 's rationale. The adjudicator is blind to writer identity and cohort meta-data and runs at $t = 0.0$. Where the weighted mean treats every judgement as exchangeable evidence, the adjudicator can weigh *arguments*, discounting an agent whose rationale misreads the essay. E2 compares the two aggregation paths directly, with BDS computed from raw agent scores in both conditions. Responding to the caution of Mirza et al. [2025], deliberation is added *after* independent judgement, never during scoring.

3.4 Data

Corpus selection

The design requires a corpus with four properties: Southeast Asian learner essays identified by L1 or country; a native-English reference cohort writing under identical task conditions; a human-assigned proficiency measure; and licence terms permitting research at this scale. A systematic search of the major public learner corpora found exactly one satisfying all four. TOEFL11 [Blanchard et al., 2013] has no Southeast Asian L1 and no native cohort, ICLE v3 lacks Southeast Asian L1s, ELLIPSE lacks L1 labels, and EFCAMDAT lacks a native reference.¹ The ICNALE corpus family [Ishikawa, 2023] satisfies all four and is adopted.

The ICNALE corpus family

ICNALE collects topic-controlled English essays from college students across ten Asian countries and regions, together with a native-English reference cohort writing under identical conditions; the three modules used are summarised in Table 3.3. Critically, every written essay addresses the same two prompts under controlled length and time conditions,

¹Corpus contents verified against their official distribution pages in July 2026: TOEFL11 via the Linguistic Data Consortium (LDC2014T06), ICLE via UCLouvain, ELLIPSE via its public repository, EFCAMDAT via the EF–Cambridge Open Language Database.

removing the prompt-variation confound of TOEFL11’s eight-topic design, and all learners carry CEFR-linked proficiency bands (A2–B2+), enabling band-restricted learner-cohort construction.

Table 3.3: ICNALE modules used in this thesis (contents as published on the official ICNALE distribution site, March 2026).

Module	Contents	Role in this thesis
Written Essays (WE) v2.6	5,600 essays, 200–300 words, two controlled topics; ten Asian countries/regions plus native ENS cohort	Primary cohorts: SEA learners and native reference
Written Essays Plus (WEP) v0.7	2,340 essays from additional countries including Vietnam, Cambodia, Laos, Myanmar	Available validating cohort for CRI_{Mekong} ; used only for demonstration essays, not scored in E1–E3
Global Archives v2.1	Rating (GRA) 140 essays, each rated by 80 raters of varied L1 and occupational backgrounds (11,200 ratings)	Multi-rater human ground truth for BDS validation (E3)

Cohort construction and sampling

Figure 3-2 summarises the sample. The Southeast Asian cohort draws from the Thailand, Indonesia, and Philippines components of WE, restricted to the two upper proficiency bands (B1_2, B2+) and stratified by topic and band. The native cohort draws from ENS, which ICNALE does not band (native writers are the corpus’s reference group). Sampling yields $n = 100$ per cohort (50 per topic) under a fixed seed. Restricting the learner cohort to its upper bands reduces how much of a cohort gap can be attributed to weak English, but it cannot equate the cohorts: any residual quality difference between upper-band learners and native writers remains inside the gap, which is therefore an unadjusted disparity (Section 3.8). A supplementary human-anchored check on the GRA essays, where human ratings exist, is reported in Section 4.4. The GRA’s 140 essays are retained in full as E3’s validation set; 13 of them (10 SEA, 3 native) also fall in the E1 sample.




Figure 3-2: Cohort construction and sampling. Essay texts remain within licensed storage; the reproducibility manifest records essay identifiers only.

Scoring target and ground truth

Every agent, including the single-agent baseline, is asked for a band from 1 to 9 on the IELTS Coherence and Cohesion descriptors (the prompts are verbatim in Appendix A). The prompts request a band from 1 to 9; Claude Sonnet and Gemini Pro returned whole bands only, while Qwen3 returned half-bands (e.g. 6.5) in 295 of its 800 scoring records and GPT-4o in 34 of its 3,700. Returned values were retained as returned, so the score resolution is one band with occasional half-bands. Panel consensus is the mean of the lens scores and so takes fractional values. The corpus’s CEFR bands are a proficiency placement (test or vocabulary-size based), not a human rating of the particular essay; association with them is therefore reported as a *proficiency-band association*, not as essay-scoring accuracy. Human ratings of the specific essays exist only for the GRA validation set. Thirteen of the 200 cohort essays also appear in the GRA set; the two analyses are reported separately. In the GRA archive each rater scores ten analytic dimensions (0–10) grouped as Language, Content, and Attitude, plus a holistic score out of 100. Per-essay rating dispersion across the 80 raters provides the human-disagreement benchmark against which BDS is compared. Two properties of the archive shape the E3 design. First, all 140 rated essays address the part-time-job prompt, so the validation set is single-topic. Second, the raters disagree substantially among themselves (mean within-essay standard deviation of the Content ratings ≈ 1.8 on the 0–10 scale). Disagreement among evaluators, human or artificial, is the nor-

mal condition of essay assessment. The question E3 asks is whether the panel’s divergence tracks the essays on which *humans* diverge.

Licence and handling

ICNALE is distributed free for research with redistribution prohibited [Ishikawa, 2023]. Essay texts are stored only in access-controlled storage and never committed to version control; the reproducibility package releases code, prompts, configurations, and essay-ID manifests only. The data are de-identified by the distributor and fall under the ethics exemption of Section 3.7.

3.5 Experimental Design

Three experiments test the four hypotheses (Table 3.4). E1 runs on all four models; E2 and E3 run on the two models with the largest E1 gaps (GPT-4o and Claude Sonnet), confining the costlier extended conditions to the evaluators where mitigation is material. All conditions share the same determinism guarantees: temperature fixed at 0.0, structured JSON outputs, exact model versions pinned in the experiment configuration, and a fixed random seed for all sampling. Every result in Chapter 4 is reproducible from the archived model outputs and the released configuration, prompts, and essay-identifier manifest; regenerating the outputs through provider APIs is best-effort (Appendix B).

Table 3.4: Overview of the three experiments: data, conditions, and the hypotheses each tests.

Exp.	Data	Conditions	Tests
E1	SEA + native cohorts ($n = 100$ each)	Single-agent baseline vs. CAMPS $k=3$, four models	H1; H2 (part)
E2	Same cohorts	CAMPS $k=5$ with and without meta-adjudicator; consensus-weight sweep ($\Delta w = 0.05$); ablation over $k \in \{1, 3, 5\}$	H2; H4
E3	GRA validation set (140 essays $\times$ 80 raters)	CAMPS $k=5$; ROC calibration of τ on a calibration split, evaluation on a hold-out split	H3

Experiment E1 (scaled replication) scores both cohorts under the single-agent baseline and the pilot’s three-agent panel on all four models, establishing whether the gap exists at scale (H1) and providing the $k=3$ reference for E2.

Experiment E2 (extended architecture) scores the same essays under the five-agent panel, with and without the meta-adjudicator. It also sweeps the consensus weights in steps of 0.05 under $\sum_i w_i = 1$, recording the cohort gap and Spearman ρ against CEFR bands at every vector; the full non-dominated set, not a single optimum, is reported and H4’s boundary sought on it. The ablation over k separates the contribution of panel size from that of the adjudicator, and the individual lenses are compared with the panel so that what the panel adds beyond its best single lens is explicit.

Experiment E3 (detection validation) computes BDS for the 140 GRA essays. Human rating dispersion provides the reference notion of a contested essay, operationalised as the top quartile of per-essay standard deviation across the 80 raters’ Content ratings, where each rater’s Content score is the mean of the four Content dimensions and dispersion is the standard deviation of those per-rater scores across the 80 raters. This is a proxy: the GRA raters disagree for many reasons, of which cultural contestation is one, so a detector of cultural divergence is here benchmarked against general disagreement. The 90%/5% operating point in H3 was fixed before E3 as the performance level the pilot’s hand-set threshold had reported (96.7% sensitivity), with 5% a conventional false-alarm tolerance. Two thresholds are calibrated on a random half of the GRA set and locked before evaluation on the held-out half: Youden’s J , and the constrained rule that H3 actually specifies (the most sensitive threshold whose calibration false-positive rate is at most 5%). Sensitivity and false-positive rate with their counts, and the area under the ROC curve, are reported.

3.6 Alternatives Considered

Training-time cultural adaptation, as in CultureLLM [Li et al., 2024], is rejected on deployment grounds: it requires weight access or fine-tuning budgets unavailable to most educational institutions, must be repeated for every model generation, and cannot be audited by the deploying institution. *Activation steering*, as in FairSteer [Li et al., 2025], is infeasible

for API-only access by construction (Section 2.2).

Single-prompt debiasing, instructing one evaluator to “score fairly across cultures”, collapses the plurality of standards into a single opaque judgement, leaving no divergence signal to calibrate a review trigger against; this instruction was not separately run. *Inter-agent debate*, as in MACD [Tan et al., 2026], is rejected because interaction can amplify bias [Mirza et al., 2025] and deliberation before scoring destroys the independence that makes BDS meaningful; CAMPS confines deliberation to the optional post-hoc adjudicator. *Recruiting external human raters* was rejected on ethics-scope and timeline grounds (Section 3.7); the GRA archive supplies eighty raters per essay without new data collection.

3.7 Ethical Considerations

This research analyses de-identified secondary data only, under an RMIT human research ethics exemption (Project ID 31036). No writer is re-identified, no essay text is redistributed, storage follows the project’s data management plan, and no participants are recruited; any extension involving human participants would require a new ethics assessment. Two considerations shape the design. First, cultural essentialism: agents are calibrated to *documented rhetorical traditions*, not nationalities, and L1 cohorts are treated as proxies for rhetorical tradition rather than identity. Second, downstream welfare: CAMPS is human-in-the-loop because the cost of a misclassified essay falls on a student. Generative AI use is disclosed in accordance with course policy.

3.8 Metrics and Statistical Protocol

Primary metric. The primary fairness quantity is the *cohort score gap*: the difference in mean panel score between the native and Southeast Asian cohorts,

$$\Delta = \overline{S^*}_{\text{native}} - \overline{S^*}_{\text{SEA}}, \quad (3.6)$$

computed on an upper-band learner cohort and the unbanded native reference. Δ is an *unadjusted scoring disparity*. On a common scale it decomposes as the human-reference cohort gap plus the difference in model error between cohorts; because the cohorts were not matched on independently assessed essay quality, Δ may include genuine quality differences as well as differential model error, and it does not by itself isolate the cultural mechanism of Definition 3.1. The treatment effects of interest are *changes* in Δ between conditions on the same essays; a fixed quality difference between cohorts enters these only through changes in how a condition maps quality to score (for example, scale compression, visible in the per-cohort standard deviations of Table 4.1). Association with the learner cohort’s proficiency bands is reported secondarily as Spearman’s ρ between the panel score and the band ordinal on the Southeast Asian cohort; a kappa statistic was considered and rejected because the 1–9 score scale and the two-level band labels share no common category system. A supplementary human-anchored check on the GRA set, added after the main analyses, is reported in Section 4.4.

Uncertainty. Every statistic carries its n . ICNALE writers contribute up to two essays (one per topic): the 200 cohort essays come from 161 writers, 39 of whom contribute both. The writer is therefore the unit of independence, and every cohort gap and gap change is reported with a 95% percentile bootstrap confidence interval that resamples *writers* (10,000 resamples, fixed seed), preserving the paired conditions within each essay. Band-association and detection statistics are reported as point estimates with their n and counts. Confidence intervals are the primary inferential device. p -values are reported only for the prespecified paired-bootstrap tests of gap change (one-sided probability of no reduction); the tests and seed were fixed in this chapter before the experiments ran but were not formally registered.

Chapter summary

CAMPS scores each essay with several culturally-calibrated lenses, averages them, and measures their disagreement. E1 tests whether the gap exists across four models (H1), E2 whether a larger panel and adjudicator help and what they cost (H2, H4), E3 whether disagreement flags contested essays (H3).

Chapter 4

Results and Discussion

This chapter reports the three experiments in the order of the hypotheses they test, then examines uncertainty and threats to validity, and closes with a discussion of what the results mean, why they arose, and how they compare with prior work. Every statistic is generated by script from the logged runs; every cohort gap and gap change carries its n and a 95% writer-cluster bootstrap confidence interval (161 writers behind the 200 essays), and band-association and detection statistics are reported as point estimates with their n and counts. Throughout, the cohort gap is an *unadjusted scoring disparity*: it may include genuine quality differences between the cohorts as well as differential model error (Section 3.8).

4.1 Experiment 1: The Cohort Gap at Scale (H1; H2 in part)

Experiment E1 scored all 200 cohort essays under the single-agent baseline and the three-agent panel on four models; the two conditions together comprised 800 scoring calls per model, all deterministic and logged. Table 4.1 reports the mean and spread of scores per cohort under each condition, the native—SEA score gap with bootstrap confidence intervals, and the gap in standard-deviation units.

Three findings emerge. First, a positive cohort gap exists on every model tested. Under the baseline evaluator, native-cohort essays received higher mean scores than upper-band

Southeast Asian essays on all four models, and each gap’s 95% confidence interval excludes zero. The gaps were 1.18 bands [0.83, 1.54] for Claude Sonnet, 0.80 [0.59, 1.02] for GPT-4o, 0.23 [0.06, 0.40] for Qwen3 14B, and 0.19 [0.06, 0.31] for Gemini Pro ($n = 100$ per cohort throughout). H1 is supported in the form in which it was tested: an unadjusted native–SEA scoring disparity, present on all four model families including an open-weight model. Because the cohorts were not matched on independently assessed essay quality, these gaps do not by themselves isolate the cultural mechanism of Definition 3.1; Section 4.4 tests that mechanism where human ratings exist, between the learner traditions the GRA contains.

Second, the *magnitude* of the gap varied across models, spanning a factor of six in raw bands. Claude Sonnet and GPT-4o exhibited the largest gaps. Gemini Pro and the open-weight Qwen3 showed gaps three to six times smaller, both scoring within a compressed range (Gemini’s native scores had a standard deviation of 0.20 bands). In pooled-SD units the spread is narrower, about 2.8-fold ($d = 1.09$ GPT-4o, 1.01 Claude, 0.43 Gemini, 0.39 Qwen3), so part of the raw heterogeneity is a difference in how much of the scale each model uses. This ordering is only partly consistent with the working theory’s claim that the filter reflects the training distribution rather than model quality, since the smallest gap belongs to a Western-trained model. Provenance is also confounded with model scale and score range in this sample, so the ordering is reported as an observation rather than a causal claim.

Third, the panel’s effect was model-dependent. Using a writer-cluster paired bootstrap (10,000 resamples), the three-agent panel reduced the gap by 0.27 bands [0.07, 0.49], or 23% ($p = 0.005$), on Claude Sonnet, and by 0.23 bands [0.12, 0.33], or 28% ($p < 0.0001$), on GPT-4o. On Qwen3 14B the change was -0.04 bands $[-0.16, 0.08]$, no clear evidence of an effect. On Gemini Pro the panel *amplified* the gap, from 0.19 to 0.45 bands (change $-0.26 [-0.39, -0.13]$). The lens-level scores locate the mechanism. Within Gemini’s panel the neutral lens scored far higher than the baseline evaluator and with a much larger cohort gap (7.18 native vs. 6.33 SEA, a 0.85-band gap on its own; $n = 100$ per cohort), while the Western and Southeast Asian lenses remained compressed. On this model the lens prompts shifted how much of the score range was used, not merely which rhetorical features were rewarded, and the panel mean inherited the expanded lens’s gap. This is the behaviour the

MALIBU benchmark warns of for persona-conditioned judging. Gemini was not carried into E2, so it bears on H2 only as a boundary condition: the panel is not a safe default on a model whose baseline gap is already small.

Figure 4-1 visualises the primary result: the gap on every model under both conditions, with its confidence intervals.




Figure 4-1: Native–SEA score gap by model and condition, with 95% writer-cluster bootstrap confidence intervals ($n=100$ per cohort). Every baseline gap excludes zero; the panel narrows the gap on the two models where it is largest, leaves it unchanged on Qwen3, and amplifies it on Gemini Pro.

The pattern in these results anticipates the fuller treatment in Sections 4.2 and 4.6: the panel’s effect differed by model, with the two largest reductions on the two models with the largest baseline gaps, no clear effect on Qwen3, and amplification on Gemini Pro. Cultural plurality in the evaluator is therefore not an intervention to be applied without a baseline audit of the kind E1 provides. Figure 4-2 summarises the four model-level observations; because the baseline gap enters both axes, and only four models were tested, the figure is descriptive and no dose–response mechanism is inferred from it.

4.2 Experiment 2: Extended CAMPS Panel (H2, H4)

Experiment E2 extended the panel from three lenses to five, adding the East Asian and Mekong lenses, and added the meta-adjudicator. It scored the same 200 cohort essays on




Figure 4-2: Gap reduction under the three-lens panel against the baseline gap, one point per model, with writer-cluster paired-bootstrap 95% confidence intervals. The panel reduced the gap on the two models with the largest baseline gaps, had no clear effect on Qwen3, and amplified the gap on Gemini Pro. Descriptive: four observations, with the baseline gap on both axes.

Table 4.1: Experiment E1: mean consensus score (SD) by cohort, $n = 100$ per cohort, and the native—SEA score gap (upper-band SEA cohort vs unbanded native reference) with 95% writer-cluster bootstrap confidence intervals (10,000 resamples); d is the gap in pooled-SD units.

Model	Condition	Mean _{SEA} (SD)	Mean _{native} (SD)	Gap [95% CI]	d
Claude Sonnet	Baseline	4.44 (1.39)	5.62 (0.91)	1.18 [0.83, 1.54]	1.01
Claude Sonnet	CAMPS $k=3$	3.73 (1.14)	4.64 (1.04)	0.91 [0.58, 1.23]	0.83
Gemini Pro	Baseline	4.83 (0.59)	5.02 (0.20)	0.19 [0.06, 0.31]	0.43
Gemini Pro	CAMPS $k=3$	5.24 (0.62)	5.69 (0.49)	0.45 [0.29, 0.61]	0.81
GPT-4o	Baseline	5.03 (0.81)	5.83 (0.65)	0.80 [0.59, 1.02]	1.09
GPT-4o	CAMPS $k=3$	4.70 (0.78)	5.27 (0.71)	0.57 [0.36, 0.79]	0.77
Qwen3 14B	Baseline	5.64 (0.59)	5.88 (0.59)	0.23 [0.06, 0.40]	0.39
Qwen3 14B	CAMPS $k=3$	5.60 (0.46)	5.87 (0.50)	0.27 [0.13, 0.40]	0.56

GPT-4o and Claude Sonnet. The cache is keyed by prompt, not by condition, so the five-lens condition reused the three E1 lens scores and added two lenses, and the adjudicated condition reused all five and added the adjudicator: 600 fresh API calls per model (400 lens, 200 adjudicator) behind 2,200 scoring records. Table 4.2 and Figure 4-4 report the ablation.

The extension did not improve mitigation. On GPT-4o, the five-lens gap was 0.55 bands [0.36, 0.75] against 0.57 [0.36, 0.79] for three lenses; the writer-cluster paired bootstrap put

the additional reduction at 0.02 bands $[-0.02, 0.06]$ ($n = 100$ per cohort, $p = 0.14$), no clear evidence of further improvement. The adjudicator behaved similarly: its final scores gave a gap of 0.52 [0.31, 0.73], a further paired change of 0.03 $[-0.03, 0.09]$ ($p = 0.16$). The panel's effect against the single-agent baseline remained robust, 0.25 bands [0.15, 0.35], or 31% ($p < 0.0001$), against 28% at $k = 3$; the 3-point difference is the 0.02-band change whose interval includes zero. What E2 rules out is not the panel but the premise that its power grows with its size. On Claude Sonnet the five-lens panel was *worse* than three lenses. The gap rose from 0.91 [0.58, 1.23] to 1.01 [0.68, 1.33], a paired change of -0.10 bands $[-0.15, -0.06]$ whose interval excludes zero, and its advantage over baseline weakened to 0.17 $[-0.03, 0.37]$. The adjudicator changed nothing on this model (gap 1.01 [0.66, 1.35], paired change 0.002 $[-0.08, 0.08]$). H2's second clause, that the panel should not degrade agreement with the corpus's proficiency bands, held on one model only. Spearman's ρ between the panel score and the CEFR band on the Southeast Asian cohort ($n = 100$; two band levels) rose on GPT-4o from 0.32 (baseline) to 0.38 ($k = 3$), 0.39 ($k = 5$) and 0.40 (adjudicated), but fell on Claude Sonnet from 0.40 to 0.34, 0.35 and 0.29. Hypothesis H2, which predicted stronger mitigation at $k = 5$ than $k = 3$, is therefore rejected on both models tested, and its band-association clause fails on Claude.

The lens-level scores suggest why, and Figure 4-3 makes the pattern visible at a glance. Each axis is one cultural lens. The native cohort's profile (outer polygon) encloses the Southeast Asian cohort's on every lens of both models: no single lens closes the gap. The shape of the two profiles, however, is nearly parallel on the two added lenses. East Asian and Mekong scored the cohorts almost in lock-step, so their inclusion diluted, rather than reinforced, the corrective contrast that the Southeast Asian lens contributes. On Claude every panel configuration was harsher than the baseline, visible as the smaller polygons: cohort means fell from 4.44 (SEA) and 5.62 (native) at baseline to 3.73 and 4.64 under three lenses, and 3.81 and 4.82 under five ($n = 100$ each), so the five-lens panel was slightly less harsh than the three-lens panel but well below the baseline. Mitigation, on this evidence, is driven by lens *composition* rather than lens *count*.

What the panel adds beyond its single most favourable lens is the sharper question, and the logged lens scores answer it. The Southeast Asian lens on its own produced a smaller

gap than any panel: 0.32 bands on GPT-4o against 0.57 for the three-lens and 0.55 for the five-lens panel, and 0.79 on Claude Sonnet against 0.91 and 1.01. The Western lens alone gave 0.54 and 0.88; the two added lenses, East Asian and Mekong, 0.51 and 0.52 on GPT-4o but 1.19 and 1.14 on Claude Sonnet, which is why the five-lens panel regressed there; and the panel’s Neutral lens the widest gaps of the pilot lenses, 0.86 and 1.06, wider than the single-evaluator baseline on GPT-4o. In gap terms, then, the panel’s reduction is the Southeast Asian lens pulling the average down, not an emergent property of combining perspectives. The Neutral lens, the panel’s regulatory ideal (Section 1.3), was not neutral in effect: asking a model to judge “argumentative merit alone” left its default expectations in place and widened the gap. What the panel adds beyond its best lens is therefore not a smaller gap but a divergence signal, whose value E3 tests.




Figure 4-3: Mean lens-level score by cohort on the five-lens panel. The native profile (red) encloses the Southeast Asian profile (blue) on every lens of both models; the added East Asian and Mekong lenses track the cohorts in parallel, explaining why enlarging the panel did not enlarge the correction.

The weight sweep addresses H4 as an exploratory analysis. Recomputing the five-lens consensus at every weight vector on a $\Delta = 0.05$ simplex grid (10,626 vectors, zero additional API calls) traces the relationship in Figure 4-5 between the absolute cohort gap and proficiency-band association (Spearman’s ρ on the banded Southeast Asian cohort). No trade-off was observed on the sample. On GPT-4o the equal-weight panel (gap 0.55, $\rho = 0.39$) was dominated by 2,594 of the weight vectors, which achieved both a smaller gap and higher association; the gap-minimising vector, which loads entirely on the South-




Figure 4-4: Cohort score gap by panel configuration and model, with 95% writer-cluster bootstrap confidence intervals. Enlarging the panel beyond $k = 3$ yields no clear further reduction on GPT-4o and a regression on Claude Sonnet whose interval excludes zero.

Table 4.2: Experiment E2: ablation over panel size k and the meta-adjudicator, with 95% writer-cluster bootstrap confidence intervals ($n=100$ per cohort).

Model	Configuration	Mean _{SEA}	Mean _{native}	Gap [95% CI]
GPT-4o	Baseline ($k=1$)	5.03	5.83	0.80 [0.59, 1.02]
GPT-4o	CAMPS $k=3$	4.70	5.27	0.57 [0.36, 0.79]
GPT-4o	CAMPS $k=5$	4.73	5.28	0.55 [0.36, 0.75]
GPT-4o	CAMPS $k=5$ + adjudicator	4.77	5.29	0.52 [0.31, 0.73]
Claude Sonnet	Baseline ($k=1$)	4.44	5.62	1.18 [0.83, 1.54]
Claude Sonnet	CAMPS $k=3$	3.73	4.64	0.91 [0.58, 1.23]
Claude Sonnet	CAMPS $k=5$	3.81	4.82	1.01 [0.68, 1.33]
Claude Sonnet	CAMPS $k=5$ + adjudicator	3.75	4.76	1.01 [0.66, 1.35]

east Asian lens, gave gap 0.32 with $\rho = 0.42$, and the non-dominated set spanned only gap 0.32–0.38 and ρ 0.42–0.43. Claude Sonnet showed the same pattern: equal weights (1.01, 0.35) dominated by 2,291 vectors, the Southeast Asian lens alone at (0.79, 0.36), and a non-dominated set spanning gap 0.79–0.90 and ρ 0.355–0.361. Three qualifications bound this reading. The gap-minimising vector being a single lens is a consequence of the aggregation rule (Section 3.2), not evidence for combining perspectives; the association measure is a two-level proficiency placement, not essay-level human quality, so “accuracy” here is a weak proxy; and the sweep optimises on the observed sample, so the non-dominated set is an in-sample description rather than a performance guarantee. Within those limits, H4’s predicted boundary was not located: on the proxy measures available, reducing the gap did

not cost association.




Figure 4-5: Weight sweep for GPT-4o: each point is one weight vector over the five lenses, plotted by absolute cohort gap and Spearman association with the CEFR band (SEA cohort). The equal-weight panel (star) is dominated; the non-dominated set clusters at the Southeast Asian lens. Exploratory and in-sample.

4.3 Experiment 3: BDS Validation (H3)

Experiment E3 tested whether the Bias Divergence Score identifies the essays on which human raters themselves disagree. The five-lens panel scored all 140 GRA essays on GPT-4o and Claude Sonnet (700 calls per model). Each essay carries ratings from 80 human raters on the four Content dimensions, whose per-essay standard deviation defines the human-disagreement ground truth. “Contested” is operationalised as the top dispersion quartile. Half the essays, selected by the prespecified seed, calibrated the thresholds; the other half was held out for evaluation.

H3 is not supported, on either model. The held-out half contained 18 contested and 52 uncontested essays. Under the rule H3 specifies, a threshold locked on the calibration half at a false-positive rate of at most 5%, GPT-4o’s BDS ($\tau = 0.63$) detected 1 of the 18 contested essays (sensitivity 0.056, false-positive rate 0.000), and for Claude Sonnet the only threshold meeting the constraint ($\tau = 0.80$) detected none (0/18). Under the more permissive Youden rule ($\tau = 0.40$ on both models), sensitivity reached 0.833 (15/18) but at false-positive rates

of 0.615 (32/52) on GPT-4o and 0.865 (45/52) on Claude. The held-out area under the ROC curve was 0.639 and 0.481, and the correlation between BDS and human dispersion across all 140 essays was $r = 0.23$ and $r = -0.13$. Table 4.3 and Figure 4-6 report both models. As a review trigger at the specified operating point, BDS did not work; at best it performed modestly better than chance, and on Claude no better than a coin flip.

Table 4.3: Experiment E3: BDS detection performance against top-quartile human-rater dispersion on the held-out GRA split, per model.

Quantity	GPT-4o	Claude Sonnet
Held-out contested / uncontested essays ($n=70$)	18 / 52	18 / 52
Youden threshold τ (calibration split)	0.40	0.40
sensitivity (held-out)	0.833 (15/18)	0.833 (15/18)
false-positive rate (held-out)	0.615 (32/52)	0.865 (45/52)
Constrained τ (calibration FPR $\leq 5\%$)	0.63	0.80
sensitivity (held-out)	0.056 (1/18)	0.000 (0/18)
false-positive rate (held-out)	0.000 (0/52)	0.019 (1/52)
Area under ROC curve (held-out)	0.639	0.481
Pearson r (BDS, human dispersion), all essays	0.232	-0.126




Figure 4-6: Left: ROC of BDS against top-quartile human dispersion on the held-out GRA split (both models). Right: BDS against the dispersion of 80 human raters; the discreteness of BDS under integer lens scores is visible as horizontal bands.

Two reasons explain this, one mechanical and one conceptual. The mechanical reason is resolution: band-level lens scores at $k = 5$ concentrate BDS onto seven distinct values, visible as horizontal bands in the scatter, leaving the score too coarse to rank 140 essays against a continuous dispersion measure; finer scores are a plausible remedy but were not tested. The conceptual reason is a construct mismatch that the case study of Section 4.4

makes concrete. BDS directly measures dispersion among prompt-conditioned scores; that this dispersion reflects *cultural* contestation is the framework’s proposed interpretation, not an established property. The dispersion of 80 raters, by contrast, aggregates every source of human disagreement, including severity, halo, and attention effects unrelated to culture. Moreover, the GRA essays vary over only a narrow dispersion range on a single topic. A detector of cultural contestation was evaluated against a benchmark of general disagreement and, unsurprisingly in hindsight, tracked it weakly. Two conclusions follow. First, the review-trigger claim of the working theory is unproven at this operating point. Second, validating it properly requires human annotation of *cultural* contestation specifically, which no public corpus currently provides, and finer-grained lens scores would address only the resolution half of the problem.

4.4 Uncertainty and Error Analysis

How the panel narrows the gap. The gap reductions in Section 4.1 are accompanied by a severity shift: the panel lowers mean scores for *both* cohorts on GPT-4o and Claude Sonnet, and lowers native-cohort scores more. On GPT-4o, mean SEA scores fell by 0.33 bands under the panel and native scores by 0.56 ($n = 100$ each); on Claude Sonnet, by 0.71 and 0.98. No cohort mean was raised: the gap narrows because both cohorts are marked down, the native cohort more. Whether this is fairness or levelling down is taken up in Section 4.6. On the low-gap model Qwen3 the shifts were 0.04 and 0.01, essentially unchanged.

Where the disparity concentrates. Disaggregating the baseline gap by proficiency band shows it largest for mid-proficiency writing. For B1_2 essays ($n = 74$), the baseline gap against the native cohort was 0.95 bands [0.72, 1.18] on GPT-4o and 1.50 [1.12, 1.85] on Claude Sonnet; for B2+ essays ($n = 26$), it fell to 0.37 [0.01, 0.74] and 0.27 [−0.20, 0.78]. The B2+ intervals are wide, and Figure 4-7 is descriptive. Two explanations are consistent with it and the design does not separate them: the working theory’s account, that intermediate writers’ rhetoric is the most distinctly local and therefore the most filtered, and the simpler account that B2+ essays are closer in quality to native essays. The subgroup pattern is reported; the cultural explanation remains a hypothesis.




Figure 4-7: Baseline cohort gap disaggregated by the Southeast Asian writers’ proficiency band: the gap is largest for mid-proficiency (B1 upper) writing on all models. Descriptive; $n = 26$ in the B2+ group.

Panel divergence. With $k = 3$ agents scoring integer bands, BDS takes few distinct values and its distribution is coarse. More telling is its direction. The working theory expects divergence to concentrate on the filtered (Southeast Asian) essays, but mean BDS was *higher* on the native cohort: 0.51 vs. 0.39 on GPT-4o at $k = 3$ (difference 0.13 [0.07, 0.18], computed before rounding) and 0.47 vs. 0.35 at $k = 5$ (0.11 [0.07, 0.16]), with the same sign on Claude Sonnet (0.10 [0.04, 0.17] at $k = 5$). The lenses disagreed more about native writing than about the writing the panel was built to protect, which anticipates E3: whatever BDS measures, it is not concentrated where the theory locates cultural contestation.

Human-anchored check. The cohort score gap is measured against an unbanded native reference and is an unadjusted disparity: part of any native–SEA gap may be a genuine quality difference that human raters would also assign, and part may be differential model error. The GRA essays, which carry 80-rater human scores, permit a supplementary test of Definition 3.1 in its own terms, between the traditions the set contains. The five-lens panel’s consensus correlated with the human Content mean at $r = 0.91$ (GPT-4o) and $r = 0.88$ (Claude Sonnet) across all 140 essays. This is evidence of rank association, not of calibrated accuracy: a score that is uniformly too low can correlate perfectly, and standardising the two scales, as the residual analysis below does, removes exactly the level and spread differences that the severity-shift question concerns. Standardising both scales and taking

the LLM–human residual by writer group, the panel showed no cohort-specific deviation between the two learner cohorts. The Southeast Asian ($n = 44$) minus East Asian ($n = 80$) residual (the remaining 12 essays, from Hong Kong, Pakistan, and Singapore, belong to neither group) difference was $+0.06$ standard deviations (SD) $[-0.07, +0.19]$ on GPT-4o and -0.04 $[-0.21, +0.13]$ on Claude Sonnet. The Western lens alone likewise showed no learner-cohort residual ($+0.05$ $[-0.13, +0.23]$ and $+0.12$ $[-0.10, +0.35]$). The Southeast Asian lens on GPT-4o, however, was *more generous* to Southeast Asian essays than the human raters ($+0.24$ $[+0.07, +0.41]$). This is a lens-calibration lean that the consensus averages out. The GRA set contains only four native-speaker essays, too few for inference; descriptively, the panel ranked them below where the human raters did (-0.22 and -0.38 SD). Two cautions apply. First, the GRA human raters themselves ordered the groups native $>$ East Asian $>$ Southeast Asian (Content means 6.35, 5.47, 4.87 on the 0–10 scale; four native essays), suggesting that raw cohort gaps partly reflect quality differences humans also perceive. Second, the analysis was added after the main analyses, is single-topic, and does not include a native–SEA contrast, which is the contrast E1 measures. It shows that under a human anchor the panel treated the two learner traditions alike; it says nothing about whether the native–SEA gap of E1 is quality difference or differential error, which this design cannot decompose.

A concrete case. One essay makes the mechanism visible. The highest-divergence Southeast Asian essay under GPT-4o’s three-lens panel (a Philippine B1_2 essay, BDS = 0.94 at $k = 3$, 0.75 at $k = 5$) opens with a proverb, “Experience is the best teacher”, and lets its position emerge inductively. Figure 4-8 shows all 22 scores it received across four models and seven evaluator roles. The Western-calibrated lens was the lowest or joint-lowest scorer on every model. Its rationale on Claude Sonnet states the objection exactly: the essay “lacks a clearly articulated thesis statement, instead relying on a general aphorism (‘Experience is the best teacher’) before gradually narrowing focus” (band 5). The Mekong lens read the same sentence as a convention: the essay “opens with a culturally familiar proverb ... that anchors the argument implicitly before the position emerges, consistent with Mekong-region framing conventions” (band 6). The SEA lens on GPT-4o described “a coherent and cohesive structure, with the thesis emerging through contextual development ... built

inductively” (band 7). The adjudicator, reading all five rationales, ruled that the Western rationale “overstates the structural deficiency” and settled on band 6. The same rhetorical choice was thus recorded as a defect by one documented interpretation of the rubric and as a convention by another. The case illustrates how different prompt-defined evaluative expectations produce different judgements of the same rhetorical structure; it does not, without an independent human assessment of this essay, establish which judgement is the more accurate, and it is therefore an illustration of the mechanism Definition 3.1 describes rather than a demonstration of it. It is also the disagreement BDS is designed to expose; that the signal did not generalise into a reliable detector across 140 essays (E3) does not make the individual case less instructive.




Figure 4-8: One essay, every lens: bands assigned to a proverb-opening, inductively organised Philippine B1_2 essay by each evaluator role on each model (E1 and E2 logs). The Western-calibrated lens is the lowest scorer on every model; dashes mark roles not run on that model.

4.5 Threats to Validity

Construct validity. The central construct risk is the use of first-language and country cohorts as proxies for rhetorical tradition. Rhetorical conventions are distributed across writing communities, not fixed properties of individuals [Connor, 2011], so cohort-level results

must not be read as claims about any individual writer; this thesis draws cohort-level conclusions only. A second construct risk concerns the agents themselves: a prompted “cultural lens” may perform a stereotype of a tradition rather than a calibration to it, and cultural-alignment evaluations of LLMs can be unstable across elicitation choices. Three design features address this. First, each CRI prompt is grounded in named contrastive-rhetoric sources rather than improvised, the Mekong lens being the partial exception noted in Section 3.3. Second, the Western and Southeast Asian lenses are checked against their corpus cohorts. The East Asian and Mekong lenses, by contrast, were not separately validated in E1–E3, though validating cohorts exist in the corpus, and their behaviour is known only through the panel. Third, all runs are deterministic, removing sampling variation from the elicitation.

Internal validity. Scores may be sensitive to prompt wording and output format. All prompts were fixed before the experiments, reproduced verbatim in Appendix A, and held constant across cohorts; a constant prompt can still interact differently with different writing characteristics, so differential prompt–essay interactions remain possible. A second threat is construct drift: the lens prompts do not merely reinterpret an unchanged rubric but add criteria (the Western prompt asks for an opening-paragraph thesis and explicit signposting; the Southeast Asian prompt admits collectivist framing and inductive development), so lens-level score differences may reflect changed scoring expectations rather than the recovery of an implicit bias in the baseline. A control panel of non-cultural prompt variants would separate these explanations and was not run (Section 5.3). Every model output was archived, so all analyses are exactly reproducible from the archive; exact regeneration through provider APIs is not guaranteed even at temperature 0. Model identifiers and run dates are recorded in Appendix B.

External validity. ICNALE’s topic control is a strength for controlled comparison and a limit on generality: two essay prompts in the main cohorts, and a single prompt in the GRA validation set. Findings therefore generalise to short argumentative learner essays of this kind; extension to other genres, lengths, and scoring rubrics is future work. The model set spans three closed systems and one open-weight model, but all are contemporary instruction-tuned LLMs; conclusions are bounded to this model class.

Statistical validity. The GRA validation set is $n = 140$, which limits the precision of E3’s sensitivity and false-positive estimates; all E3 statistics are reported with their n , and the calibration/evaluation split is fixed by seed before analysis. Across all experiments, confidence intervals are the primary inferential device and no post-hoc significance testing is performed (Section 3.8).

4.6 Discussion

In plain terms, the experiments found the following. On every model tested, current language models scored the sampled native-English essays higher than the sampled upper-band Southeast Asian essays, by a margin that differed six-fold across models in raw bands and about three-fold in standardised units. Because the two groups were not matched on independently judged quality, this is an unadjusted disparity rather than a measured cultural bias. A small panel of culturally framed evaluators reduced the disparity by about a quarter on the two models where it was largest, entirely by marking both cohorts down (the native cohort more), and amplified it on the model where it was smallest. Enlarging the panel from three lenses to five added nothing and on one model subtracted; the reduction was no larger than that of the Southeast Asian lens alone. Reweighting found no trade-off between the gap and proficiency-band association on the sample. And the panel’s internal disagreement did not identify the essays human raters dispute.

Set against prior findings, the results replicate in one respect and fail to generalise or transfer in two. The Part A pilot reported a 20.1% gap reduction under the three-agent panel on 30 essays; the 23–28% reductions observed here on 200 essays replicate that effect on two of four model families, but it did not extend to the other two (null on Qwen3, amplification on Gemini). The pilot also reported 96.7% BDS sensitivity under a hand-set threshold at $k = 3$; the five-lens BDS, ROC-calibrated against 80-rater ground truth on 140 essays, reached 0.056 sensitivity at the specified false-positive rate, consistent with a small- n artefact in the pilot (the $k = 3$ signal itself was not validated on the GRA). The existence finding accords with Gayed [2026], who documented L1-linked scoring variation on TOEFL11, and adds a native reference and a cross-family comparison that single-model studies cannot

see. The mitigation finding accords only superficially with the panel-of-evaluators result of Verga et al. [2024], that judge diversity reduces judge bias: on this evidence the reduction is attributable to one lens, not to diversity as such. The panel-size finding shows that the accuracy gains CAFES and RES report from adding agents [Su et al., 2025, Jang et al., 2025] do not extend to fairness on this evidence. Finally, the Gemini gap amplification is the behaviour the MALIBU benchmark predicted for persona-conditioned judging [Mirza et al., 2025], and the Neutral lens’s expanded score range on that model is consistent with the elicitation instability reported by Khan et al. [2025].

Three aspects of these results bear on the working theory. First, a gap on all four model families, including an open-weight model, is what Definition 3.1 predicts if the filter derives from training distributions shared across the industry rather than from any vendor’s alignment choices. Second, the cross-model ordering is only partly consistent with that account, since the smallest gap belongs to a Western-trained model (Section 4.1). Third, the lens-level results temper the theory’s account of *why* the panel helps. The reduction on GPT-4o and Claude was carried by the Southeast Asian lens, which alone produced a smaller gap than any panel, while the Neutral prompt produced the widest gap; since the lens prompts add scoring criteria rather than merely reinterpreting a fixed rubric (Section 4.5), the most economical reading is that culturally framed prompts *change the scoring expectations applied*, and that the framing favourable to Southeast Asian rhetoric narrows the disparity. Whether that narrowing recovers quality the baseline overlooked, or simply relaxes a criterion, is not settled by this design. Gemini Pro’s amplified gap shows that the same prompts can move a weakly disparate evaluator the other way. CAMPS should therefore be deployed only after a baseline audit of the kind E1 provides.

A natural objection deserves direct treatment: are current frontier models not simply too capable for cultural bias to matter? The two largest gaps belonged to the two frontier closed models in the set, which is at least consistent with capability and evaluative neutrality being different properties; the ordering is confounded with scale and scoring style, so no more is claimed. What the variation across vendors does establish is that the choice of scoring model is itself a fairness decision, one that institutions currently make without evidence. Even under a future model with no measurable gap, the part of this thesis that would stand

is the audit: band-restricted cohort measurement with writer-cluster confidence intervals, re-runnable on any model from the released manifest. The divergence-based trigger is not yet part of that instrument, because E3 did not validate it.

The severity shift documented in Section 4.4 raises a question the fairness literature knows as levelling down: the panel achieves its narrowing by scoring native essays lower, not by treating Southeast Asian essays more generously. If the baseline native scores were inflated by the same familiarity that penalised Southeast Asian essays, this is a correction; if not, it is a cost. The evidence gathered here does not settle the question. Proficiency-band association is available only for the learner cohort, so it says nothing about whether the native scores removed were informative; and E3’s construct mismatch means the GRA human ratings cannot arbitrate. Settling correction-versus-cost requires graded human ground truth on the cohort essays themselves, which is future work.

The E2 and E3 results together discipline how the framework should be described. What survives is CAMPS as an audit instrument that also delivers a model-dependent reduction of an unadjusted disparity, with no observed trade-off against proficiency-band association. What does not yet survive is the panel as an autonomous triage instrument: BDS at band-level resolution did not meet its specified operating point, and its dispersion was concentrated on the native rather than the Southeast Asian essays. H3 tested a prespecified hypothesis on its held-out split and carries the same evidentiary weight as the H1 confirmation; H4 is exploratory.

Chapter summary

An unadjusted native–SEA disparity exists on every model tested and varies six-fold between them (H1). A three-lens panel reduces it by about a quarter on the two largest-gap models, no more than its Southeast Asian lens alone, has no clear effect on Qwen3, and amplifies it on Gemini Pro; five lenses add nothing on GPT-4o and widen the gap on Claude, and the adjudicator adds nothing (H2). Reweighting showed no gap–band-association trade-off (H4, exploratory). Panel disagreement did not identify contested essays at the specified operating point (H3).

Chapter 5

Conclusion

5.1 Conclusions Against the Hypotheses

This thesis asked whether LLM essay evaluators penalise valid non-Western rhetoric, and whether a culturally plural panel can mitigate and detect that penalty. Table 5.1 answers each hypothesis against the evidence of Chapter 4.

The table is read as follows. An unadjusted native–SEA scoring disparity exists on every model family tested, six-fold heterogeneous in magnitude and larger for the intermediate-band essays; because the cohorts were not matched on independently assessed quality, this disparity is consistent with evaluative epistemic filtering but does not establish it (H1). Culturally framed scoring prompts changed the disparity in a model-dependent way, reducing it on the two models where it was largest, by marking both cohorts down (the native cohort more) and by no more than the Southeast Asian lens alone achieved, and amplifying it on the model where it was smallest; five lenses added nothing on GPT-4o and widened the gap on Claude, and the adjudicator added nothing (H2). H4 could not be tested against graded ground truth; on the proficiency-band proxy, no trade-off was observed (H4). The panel’s divergence signal did not meet its specified operating point, and was larger on native than on Southeast Asian essays, so the review-trigger use of BDS remains a hypothesis rather than a validated capability (H3).

Turning to the research questions, RQ1 is answered in part: a scoring disparity against Southeast Asian writing is present and general, its strength is a property of the individual

Table 5.1: Verdicts on the four hypotheses. All statistics from Chapter 4; $n=100$ per cohort (E1/E2), $n=70$ held-out (E3); 95% writer-cluster bootstrap CIs. The cohort gap is an unadjusted disparity (cohorts not matched on independently judged quality).

Hypothesis	Verdict	Decisive evidence
H1 existence & generality	Supported as an unadjusted disparity	Native–SEA gap positive on all four families, every CI excluding zero: 1.18 (Claude), 0.80 (GPT-4o), 0.23 (Qwen3), 0.19 (Gemini) bands ($d = 1.01, 1.09, 0.39, 0.43$); larger for B1_2 essays. Cultural mechanism not isolated by the design
H2 mitigation, $k=5 > k=3$	Not supported (GPT-4o); rejected (Claude)	$k=3$ panel cut the gap 23–28% on the two largest-gap models, no more than its SEA lens alone; $k=5$ added 0.02 $[-0.02, 0.06]$ on GPT-4o, regressed $-0.10 [-0.15, -0.06]$ on Claude; adjudicator null on both; band association fell on Claude. Scope note: the $k=3$ panel amplified the gap on Gemini ($0.19 \rightarrow 0.45$), not carried into E2
H3 detection $\geq 90\%/\leq 5\%$	Not supported	At the specified $FPR \leq 5\%$: held-out sensitivity 0.056 (GPT-4o) and 0.000 (Claude). Held-out AUC 0.639 and 0.481; $r(\text{BDS, human dispersion}) = 0.23$ and -0.13
H4 trade-off boundary	Not testable against graded ground truth; no trade-off on the proxy (exploratory)	Graded ground truth is absent for the cohort essays, so accuracy is proxied by band association. 10,626-vector sweep: equal-weight panel dominated on both models; gap $0.55 \rightarrow 0.32$ with association $0.39 \rightarrow 0.42$ (GPT-4o), $1.01 \rightarrow 0.79$ with $0.35 \rightarrow 0.36$ (Claude); minimum-gap vector is the SEA lens alone, by construction

model, and isolating its cultural cause requires matched human assessments this design did not include. RQ2 is answered conditionally. Culturally framed prompts can narrow the disparity at inference time on the models where it is material, and panel scores track 80-rater human consensus in rank ($r \approx 0.9$; calibration was not assessed); but the narrowing has not been shown to be an improvement in fairness rather than a change of scoring criteria, and the self-diagnostic signal is not yet reliable enough to route essays autonomously.

5.2 Contributions

This thesis makes five contributions. Conceptually, it formalises *evaluative epistemic filtering* as a distinct bias mechanism in LLM-driven scoring (Definition 3.1), grounded in epis-

temic injustice and contrastive rhetoric. Architecturally, it specifies CAMPS, a culturally-calibrated scoring panel with a divergence-based human-review trigger, extending multi-perspective bias mitigation from generation [Dorleon and Shujaa, 2026] to evaluation. Empirically, it evaluates the framework on topic-controlled ICNALE cohorts across four model families, establishing the disparity at scale and characterising where the panel changes it; and it disconfirms two working assumptions, that panel benefit grows with panel size and that inter-judge divergence transfers directly into a detection signal. Methodologically, it introduces ROC calibration of the Bias Divergence Score against the disagreement of eighty human raters, an evaluation design that exposed the construct gap between cultural contestation and general rater disagreement. Practically, it gives deploying institutions a re-runnable baseline audit and an account of what narrowing the disparity did and did not cost on this evidence.

5.3 Limitations

Beyond the validity threats analysed in Section 4.5, four limitations bound this work. First, corpus scope: the evidence rests on short, topic-controlled argumentative essays by university students; high-stakes examination writing differs in genre, length, and stakes. Second, the cultural lenses cover five perspectives and leave others (notably Middle Eastern, South Asian, and African traditions) to future work; ICNALE’s WEP module does contain South Asian writers, so the constraint is the absence of validating cohorts and documented rhetorical grounding for a lens, not of essays. Third, CAMPS is an inference-time mitigation: it treats the symptom at the point of judgement and cannot correct the training-distribution causes of the filter; a panel of biased judges may be fairer than one, but is not unbiased. Fourth, the framework multiplies inference cost by the panel size and routes contested essays to humans, so its operational cost scales with exactly the essays that most need care; whether that is acceptable is an institutional decision.

Stated plainly, the scope within which the findings hold is this. CAMPS narrowed the disparity on the two frontier models whose measured gaps were largest (0.8 bands or more, a description of these two models rather than a threshold), on short argumentative essays by

upper-band Southeast Asian learners, with three lenses. It had no clear effect on a model with a small gap and amplified the gap on the model with the smallest. It should not be applied without first measuring the disparity it is meant to treat. Five internal limitations qualify the numbers. The cohorts were not matched on independently assessed quality, so the cohort gap is an unadjusted disparity; the human-anchored check could compare only the learner cohorts, since the rated archive holds four native essays. The lens prompts add scoring criteria rather than reinterpreting a fixed rubric, so the panel’s effect may reflect changed expectations rather than a corrected bias; a control panel of non-cultural prompt variants, which would separate these explanations, was not run. BDS was evaluated at band-level resolution (seven distinct values), which limits any detector built on it. The “accuracy” axis of the weight sweep is association with a two-level proficiency placement, not essay-level human quality, so no claim about scoring accuracy rests on it. And the 80-rater correlation of about 0.9 is evidence of rank association on a single-topic set, not of calibration.

5.4 Future Work

The null and unresolved results define the nearest future work. First, re-validating BDS with continuous, confidence-weighted lens scores against annotations of specifically *cultural* contestation, so that the detector is tested on the construct it measures. Alongside this, the non-cultural control panel that would isolate what cultural plurality adds over prompt averaging should be run. Second, obtaining graded human ratings on the cohort essays, so that Definition 3.1 can be tested as an LLM–human residual between native and Southeast Asian writers rather than read from a raw cohort gap. Beyond these, the panel could be extended to further rhetorical traditions as validating corpora become available, including a Middle Eastern lens for which the literature is rich but public rated corpora are lacking. A theoretical treatment is a further direction: the formal guarantees established for multi-perspective generation [Dorleon and Shujaa, 2026] suggest an analogous proposition stating conditions under which panel consensus provably narrows cohort gaps, independent of any particular model generation. Finally, a classroom pilot of BDS-triggered review.

Appendix A

CRI System Prompts

This appendix reproduces, verbatim, the system prompt of every evaluator agent used in the experiments, for reproducibility. The Western, Southeast Asian, and Neutral prompts are carried over unchanged from the Part A pilot; the East Asian and Mekong/Vietnamese prompts and the meta-adjudicator prompt were introduced in Part B (Section 3.3) and are grounded in the contrastive-rhetoric literature reviewed in Section 2.5. The user turn, identical for every agent, supplies the essay text and the instruction to respond in JSON with a band score and a short rationale; it is constructed in `src/run_experiment.py` in the released repository.

Baseline (single evaluator)

You are an essay evaluator. Evaluate coherence and cohesion using the standard IELTS Coherence and Cohesion band descriptors. Score 1-9.

CRI_{Western} (Agent 1)

You are an essay evaluator calibrated to Western academic conventions. Evaluate coherence and cohesion based on: (1) clear thesis statement in the opening paragraph, (2) linear logical progression with explicit topic sentences, (3) direct evidence marshalling with assertive hedging, (4) explicit cohesive devices (however, therefore, more-

over). Score 1-9 on the IELTS Coherence and Cohesion band descriptors.

CRI_{SEA} (Agent 2)

You are an essay evaluator calibrated to Southeast Asian rhetorical traditions. Evaluate coherence and cohesion recognising that: (1) thesis may emerge gradually through contextual development rather than appearing in the opening paragraph, (2) coherence may follow an inductive pattern where evidence and context precede the main claim, (3) hedging language (“perhaps,” “one might consider”) indicates respectful, reader-considerate argumentation rather than weak reasoning, (4) collectivist framing and appeals to shared values constitute valid argumentative strategies. Score 1-9 on the IELTS Coherence and Cohesion band descriptors.

CRI_{EastAsian} (Agent 3, introduced in Part B)

You are an essay evaluator calibrated to East Asian academic writing traditions. Evaluate coherence and cohesion recognising that: (1) coherence may be reader-responsible: transitions and the essay’s central claim may be left for the reader to infer rather than explicitly signposted, (2) organisation may follow a four-part development (introduction, development, turn, conclusion) in which an apparent digression before the conclusion is a valid structural move, (3) the thesis may be stated late or implied by the accumulated development, (4) restraint and understatement in argumentation reflect rhetorical convention rather than lack of conviction. Score 1-9 on the IELTS Coherence and Cohesion band descriptors.

CRI_{Mekong} (Agent 4, introduced in Part B)

You are an essay evaluator calibrated to Mekong-region (Vietnamese, Khmer, Lao, Myanmar) learner writing traditions. Evaluate coherence and cohesion recognising that: (1) arguments may be anchored in proverbs, family or community narratives that carry the essay’s claim implicitly, (2) contextual and relational framing may precede

the position statement, (3) deference markers and indirectness reflect politeness conventions rather than uncertainty, (4) the conclusion may perform the thesis rather than restate it. Score 1-9 on the IELTS Coherence and Cohesion band descriptors.

CRI_{Neutral} (Agent 5)

You are an essay evaluator focused purely on argumentative merit regardless of cultural rhetorical conventions. Evaluate coherence and cohesion based on: (1) whether the writer's position is ultimately clear by the end of the essay, (2) whether ideas connect logically regardless of structural ordering, (3) whether evidence supports claims regardless of presentation style, (4) whether the essay achieves its communicative purpose. Score 1-9 on the IELTS Coherence and Cohesion band descriptors.

Meta-adjudicator (introduced in Part B)

You are a meta-adjudicator for an essay-scoring panel. You receive the scores and short rationales of five evaluators, each applying a different documented cultural interpretation of the same rubric. Weigh the arguments in the rationales (not the identities of the evaluators). Discount a rationale that misidentifies the essay's structure or contradicts the essay's text. You are not told anything about the writer. Respond with JSON: {"score": <band 1-9>, "confidence": <0-1>, "rationale": "<3-4 sentences>"}.

Appendix B

Reproducibility

All experiments are reproducible from the public code repository (GitHub: [mikenk2010/vs4139514-Final-Thesis-CAMPS-Code](https://github.com/mikenk2010/vs4139514-Final-Thesis-CAMPS-Code)), which contains the pipeline, prompts, configurations, analysis code, and essay-identifier manifests. Corpus texts are not distributed, per the ICNALE licence; they are obtainable free of charge through the ICNALE registration process.

Pipeline

Experiments are driven by a single configuration file (`configs/experiment.yaml`) specifying models, conditions, cohorts, and the sampling seed (4139514). Every scoring call uses temperature 0.0 with structured JSON output, is cached on disk keyed by (model, agent, essay, prompt), and is logged as one JSON line with timestamp, model identifier, and token usage; runs are resumable. Two senses of reproducibility are distinguished. Every table and figure in Chapter 4 is exactly reproducible from the archived model outputs and manifests. Regenerating the outputs themselves through provider APIs is not guaranteed to be bit-identical: temperature 0 reduces but does not eliminate generation variability, and providers document reproducibility as best-effort. Scoring ran between 18 August and 2 September 2026. Each run record counts calls issued and calls skipped on resume; because the cache key excludes the experimental condition, lens prompts shared between conditions (and the 13 GRA essays also in the E1 sample) were served from the cache, so

the 5,971 cache entries are the complete set of distinct API responses, including the 249 archived Gemini 2.5 calls. The sampling manifest (essay identifiers only) is generated by `src/sample_manifest.py`; the GRA validation manifest and the human-dispersion ground truth are generated by `src/make_gra_manifest.py` from the archive’s rating workbook.

Analysis

All results tables and figures in Chapter 4 are generated by `src/make_tables.py` and `src/sweep_weights.py` directly from the experiment logs; no statistic is computed or transcribed by hand. Cohort gaps, writer-cluster bootstrap confidence intervals (10,000 resamples, fixed seed), Spearman band association, BDS, and ROC calibration are implemented in `src/metrics.py` with unit tests.

Model versions

The exact model identifiers requested through each provider’s API, pinned in the experiment configuration (`configs/experiment.yaml`) and recorded in every log entry, were: `gpt-4o` (OpenAI); `claude-sonnet-4-6` (Anthropic); `qwen3:14b` (Alibaba, open weights, executed locally via Ollama); and Google’s `gemini-3.1-pro-preview`. One model change occurred mid-study. The study began with `gemini-2.5-pro`, which the provider withdrew from new API access partway through data collection. The 249 calls completed against it were archived and excluded from all analyses (mixing scores from two model versions within one experimental cell would confound the comparison), and the full Gemini condition was rerun from scratch on `gemini-3.1-pro-preview`. The archived log is retained in the repository for audit. The local model ran under Ollama with its default quantisation for the `qwen3:14b` tag; provider models are dated snapshots controlled by the provider, and the exact revision served is not exposed for all of them, so the identifiers above are the requested identifiers rather than a complete frozen environment. Client library versions are pinned in the repository. The sampling manifest was frozen with seed 4139514 before any scoring began. Human dispersion for E3 was computed per essay as the standard deviation, across

the 80 raters, of each rater's mean of the four Content dimensions; raters with a missing Content rating for an essay were excluded for that essay.

Bibliography

- Daniel Blanchard, Joel Tetreault, Derrick Higgins, Aoife Cahill, and Martin Chodorow. TOEFL11: A corpus of non-native English. Technical Report RR-13-24, Educational Testing Service, 2013. Corpus distributed as LDC2014T06, Linguistic Data Consortium.
- Xiaoqing Cheng, Ruizhe Chen, Hongying Zan, Yuxiang Jia, and Min Peng. BiasFilter: An inference-time debiasing framework for large language models. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, 2025.
- Ulla Connor. *Intercultural Rhetoric in the Writing Classroom*. University of Michigan Press, Ann Arbor, MI, 2011.
- Ginel Dorleon and Shirin Shujaa. Uncovering cultural biases and stereotypes in large language models. In *Proceedings of the 27th International Conference on Asia-Pacific Digital Libraries (ICADL)*, volume 16242 of *Lecture Notes in Computer Science*, pages 64–77. Springer, 2026. doi: 10.1007/978-981-95-4861-3_5.
- Ginel Dorleon, Imen Megdiche, Nathalie Bricon-Souf, and Olivier Teste. Feature selection under fairness constraints. In *Proceedings of the 37th ACM/SIGAPP Symposium on Applied Computing (SAC '22)*, pages 1125–1127. ACM, 2022. doi: 10.1145/3477314.3507168.
- Ginel Dorleon, Imen Megdiche, Nathalie Bricon-Souf, et al. FAPFID: A fairness-aware approach for protected features and imbalanced data. In *Transactions on Large-Scale Data- and Knowledge-Centered Systems*, volume 13840 of *Lecture Notes in Computer Science*, pages 107–125. Springer, 2023. doi: 10.1007/978-3-662-66863-4_5.
- Miranda Fricker. *Epistemic Injustice: Power and the Ethics of Knowing*. Oxford University Press, Oxford, 2007.
- Isabel O. Gallegos, Ryan A. Rossi, Joe Barrow, Md Mehrab Tanjim, Sungchul Kim, Franck Dernoncourt, Tong Yu, Ruiyi Zhang, and Nesreen K. Ahmed. Bias and fairness in large language models: A survey. *Computational Linguistics*, 50(3):1097–1179, 2024.
- John Maurice Gayed. Investigating first-language bias in LLM-based automated essay scoring: A cross-prompt evaluation of an open-weight AI model on TOEFL essays. arXiv preprint, 2026. arXiv:2607.14605.

- John Hinds. Reader versus writer responsibility: A new typology. In Ulla Connor and Robert B. Kaplan, editors, *Writing Across Languages: Analysis of L2 Text*, pages 141–152. Addison-Wesley, Reading, MA, 1987.
- Shin’ichiro Ishikawa. *The ICNALE Guide: An Introduction to a Learner Corpus Study on Asian Learners’ L2 English*. Routledge, 2023. ISBN 9781032180250.
- Jinhee Jang, Ayoung Moon, Minkyung Jung, YoungBin Kim, and Seung Jin Lee. LLM agents at the roundtable: A multi-perspective and dialectical reasoning framework for essay scoring. arXiv preprint, 2025. arXiv:2509.14834.
- Robert B. Kaplan. Cultural thought patterns in inter-cultural education. *Language Learning*, 16(1–2):1–20, 1966.
- Ali Keramati, Shiyuan Zhou, Sharad Mehrotra, and Mark Warschauer. MADRAG: Multi-agent debate with retrieval-augmented generation for training-free analytic essay scoring. arXiv preprint, 2026. arXiv:2606.06754.
- Ariba Khan, Stephen Casper, and Dylan Hadfield-Menell. Randomness, not representation: The unreliability of evaluating cultural alignment in LLMs. arXiv preprint, 2025.
- Cheng Li, Mengzhuo Chen, Jindong Wang, Sunayana Sitaram, and Xing Xie. CultureLLM: Incorporating cultural differences into large language models. In *Advances in Neural Information Processing Systems*, 2024.
- Yichen Li, Zhiting Fan, Ruizhe Chen, Xiaotang Gai, Luqi Gong, Yan Zhang, and Zuozhu Liu. FairSteer: Inference time debiasing for LLMs with dynamic activation steering. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 11293–11312, Vienna, Austria, 2025. Association for Computational Linguistics. doi: 10.18653/v1/2025.findings-acl.589.
- Imran Mirza, Cole Huang, Ishwara Vasista, Asli Akalin, Sean O’Brien, Rohan Patil, and Kevin Zhu. MALIBU benchmark: Multi-agent LLM implicit bias uncovered, 2025. arXiv:2507.01019.
- Austin Pack, Alex Barrett, and Juan Escalante. Large language models and automated essay scoring of English language learner writing: Insights into validity and reliability. *Computers and Education: Artificial Intelligence*, 6:100234, 2024.
- Timo Schick, Sahana Udupa, and Hinrich Schütze. Self-diagnosis and self-debiasing: A proposal for reducing corpus-based bias in NLP. *Transactions of the Association for Computational Linguistics*, 9:1408–1424, 2021. doi: 10.1162/tac1_a_00434.
- Shirin Shujaa and Ginel Dorleon. Detecting framing bias in news via probabilistic graphical modeling. In *Proceedings of the ACM/IEEE Joint Conference on Digital Libraries (JCDL)*, pages 245–248. IEEE, 2025. doi: 10.1109/JCDL67857.2025.00037.

- Shirin Shujaa, Ginel Dorleon, and A. Tang. How LLMs handle cultural bias: Reactions to Asian minority historical narratives. In *Proceedings of the 27th International Conference on Asia-Pacific Digital Libraries (ICADL)*, volume 16242 of *Lecture Notes in Computer Science*, pages 39–52. Springer, 2026. doi: 10.1007/978-981-95-4861-3_3.
- Jiamin Su, Yibo Yan, Zhuoran Gao, Han Zhang, Xiang Liu, and Xuming Hu. CAFES: A collaborative multi-agent framework for multi-granular multimodal essay scoring. arXiv preprint, 2025. arXiv:2505.13965.
- Yosephine Susanto, Adithya Venkatadri Hulagadri, Jann Railey Montalan, Jian Gang Ngui, Xianbin Yong, Wei Qi Leong, Hamsawardhini Rengarajan, Peerat Limkonchotiwat, Yifan Mai, and William Chandra Tjhi. SEA-HELM: Southeast Asian holistic evaluation of language models. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 12308–12336, Vienna, Austria, 2025. Association for Computational Linguistics. doi: 10.18653/v1/2025.findings-acl.636.
- Qian Tan, Lei Jiang, Yuting Zeng, Shuoyang Ding, and Xiaohua Xu. Mitigating cultural bias in LLMs via multi-agent cultural debate. arXiv preprint, 2026. arXiv:2601.12091.
- Yan Tao, Olga Viberg, Ryan S. Baker, and René F. Kizilcec. Cultural bias and cultural alignment of large language models. *PNAS Nexus*, 3(9):pgae346, 2024. doi: 10.1093/pnasnexus/pgae346.
- Sowmya Vajjala. Automated assessment of non-native learner essays: Investigating the role of linguistic features. *International Journal of Artificial Intelligence in Education*, 28: 79–105, 2018.
- Pat Verga, Sebastian Hofstätter, Sophia Althammer, Yixuan Su, Aleksandra Piktus, Arkady Arkhangorodsky, Minjie Xu, Naomi White, and Patrick Lewis. Replacing judges with juries: Evaluating LLM generations with a panel of diverse models. arXiv preprint, 2024. arXiv:2404.18796.
- Yun Wang, Zhaojun Ding, Xuansheng Wu, Siyue Sun, Ninghao Liu, and Xiaoming Zhai. AutoSCORE: Enhancing automated scoring with multi-agent large language models via structured component recognition. arXiv preprint, 2025. arXiv:2509.21910.
- Kaixun Yang, Mladen Raković, Dragan Gašević, and Guanliang Chen. Does the prompt-based large language model recognize students’ demographics and introduce bias in essay scoring? In *Artificial Intelligence in Education (AIED 2025)*, volume 15878 of *Lecture Notes in Computer Science*. Springer, 2025. doi: 10.1007/978-3-031-98417-4_6.